

# road-acc-fin

March 31, 2023

## 0.1 Problem Statement

We are working with a dataset on accident severity, and our goal is to build a machine learning model that can accurately predict the severity of accidents based on a range of different features. Our main objective is to identify the most important features for predicting accident severity and to develop a model that can generalize well to new data.

The dataset contains information about road traffic accidents in the UK, from the year 2005 to 2021. The columns in the dataset contain the following information:

- `accident_index`: Unique identifier for each accident
- `accident_year`: Year in which the accident occurred
- `accident_reference`: Reference number for the accident
- `location_easting_osgr`: Easting (X-coordinate) of accident location in British National Grid
- `location_northing_osgr`: Northing (Y-coordinate) of accident location in British National Grid
- `longitude`: Longitude of accident location
- `latitude`: Latitude of accident location
- `police_force`: Police force that attended the accident
- `accident_severity`: Severity of the accident, ranging from 1 (fatal) to 3 (slight)
- `number_of_vehicles`: Number of vehicles involved in the accident
- `number_of_casualties`: Number of casualties in the accident
- `date`: Date of the accident
- `day_of_week`: Day of the week on which the accident occurred
- `time`: Time of the accident
- `local_authority_district`: Local authority district in which the accident occurred
- `local_authority_ons_district`: ONS code for the local authority district in which the accident occurred
- `local_authority_highway`: Highway authority responsible for the road on which the accident occurred
- `first_road_class`: Class of the first road involved in the accident, ranging from 1 (motorway) to 6 (other)
- `first_road_number`: Number of the first road involved in the accident
- `road_type`: Type of road on which the accident occurred, ranging from 1 (motorway) to 6 (footpath or bridleway)
- `speed_limit`: Speed limit on the road on which the accident occurred
- `junction_detail`: Type of junction at which the accident occurred
- `junction_control`: Control of the junction at which the accident occurred
- `second_road_class`: Class of the second road involved in the accident, ranging from 0 (not

- at junction or within 20 metres) to 6 (other)
- second\_road\_number: Number of the second road involved in the accident
- pedestrian\_crossing\_human\_control: Whether a pedestrian crossing with a human controller was present at the accident location
- pedestrian\_crossing\_physical\_facilities: Whether a pedestrian crossing with physical facilities was present at the accident location
- light\_conditions: Light conditions at the time of the accident
- weather\_conditions: Weather conditions at the time of the accident
- road\_surface\_conditions: Road surface conditions at the time of the accident
- special\_conditions\_at\_site: Whether any special conditions were present at the accident location
- carriageway\_hazards: Whether any carriageway hazards were present at the accident location
- urban\_or\_rural\_area: Whether the accident occurred in an urban or rural area
- did\_police\_officer\_attend\_scene\_of\_accident: Whether a police officer attended the scene of the accident
- trunk\_road\_flag: Whether the road on which the accident occurred is a trunk road
- lsoa\_of\_accident\_location: Lower Super Output Area in which the accident occurred

This information can be used to understand the factors that contribute to road traffic accidents and to develop strategies to prevent such accidents in the future.

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: !unzip -q "/content/drive/MyDrive/code_files/Sri/road_acc_4/Data.zip"
```

```
[4]: # Libraries to help with reading and manipulating data
import pandas as pd
import numpy as np

# Libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# To tune model, get different metric scores, and split data
from sklearn.metrics import (
    f1_score,
    accuracy_score,
    recall_score,
    precision_score,
    confusion_matrix,
    roc_auc_score
)
from sklearn.model_selection import train_test_split, StratifiedKFold,
↪cross_val_score
```

```

# To be used for data scaling and one hot encoding
from sklearn.preprocessing import StandardScaler, MinMaxScaler, OneHotEncoder
from sklearn import metrics

# To impute missing values
from sklearn.impute import SimpleImputer

# To oversample and undersample data
from imblearn.over_sampling import SMOTE
from imblearn.under_sampling import RandomUnderSampler

# To do hyperparameter tuning
from sklearn.model_selection import RandomizedSearchCV

# To define maximum number of columns to be displayed in a dataframe
pd.set_option("display.max_columns", None)

# To help with model building
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

# To suppress warnings
import warnings

warnings.filterwarnings("ignore")

```

```
[5]: data = pd.read_csv("/content/Data/Accidents-2021.csv")
```

```
[6]: data.head()
```

```
[6]:
```

	accident_index	accident_year	accident_reference	location_easting_osgr	\
0	2021010000000.0	2021	10287148	521508.0	
1	2021010000000.0	2021	10287149	535379.0	
2	2021010000000.0	2021	10287151	529701.0	
3	2021010000000.0	2021	10287155	525312.0	
4	2021010000000.0	2021	10287157	512144.0	

	location_northing_osgr	longitude	latitude	police_force	\
0	193079.0	-0.246102	51.623425	1	
1	180783.0	-0.050574	51.509767	1	
2	170398.0	-0.136152	51.417769	1	
3	178385.0	-0.196411	51.490536	1	
4	171526.0	-0.388169	51.431649	1	

	accident_severity	number_of_vehicles	number_of_casualties	date	\
0	3	3	1	1/1/2021	
1	2	2	3	1/1/2021	
2	2	2	4	1/1/2021	
3	1	1	1	1/1/2021	
4	3	4	1	1/1/2021	

	day_of_week	time	local_authority_district	local_authority_ons_district	\
0	6	2:05	-1	E09000003	
1	6	3:30	-1	E09000030	
2	6	4:07	-1	E09000022	
3	6	4:26	-1	E09000020	
4	6	3:10	-1	E09000018	

	local_authority_highway	first_road_class	first_road_number	road_type	\
0	E09000003	6	0	6	
1	E09000030	3	1203	3	
2	E09000022	4	272	6	
3	E09000020	3	3220	2	
4	E09000018	5	0	6	

	speed_limit	junction_detail	junction_control	second_road_class	\
0	30	9	4	6	
1	30	7	2	3	
2	30	9	2	5	
3	30	9	4	6	
4	20	3	4	6	

	second_road_number	pedestrian_crossing_human_control	\
0	0	0	
1	1204	0	
2	0	0	
3	0	0	
4	0	0	

	pedestrian_crossing_physical_facilities	light_conditions	\
0	0	4	
1	5	4	
2	5	4	
3	4	4	
4	0	4	

	weather_conditions	road_surface_conditions	special_conditions_at_site	\
0	7	4	1	
1	1	1	0	
2	1	1	0	
3	1	1	0	

4	1	1	0
---	---	---	---

	carriageway_hazards	urban_or_rural_area	\
0	0	1	
1	0	1	
2	0	1	
3	0	1	
4	0	1	

	did_police_officer_attend_scene_of_accident	trunk_road_flag	\
0	1	2	
1	1	2	
2	1	2	
3	1	2	
4	1	2	

	lsoa_of_accident_location
0	E01000263
1	E01004303
2	E01003146
3	E01002847
4	E01002608

```
[7]: data.describe().T
```

```
[7]:
```

	count	mean	\
accident_year	101087.0	2021.000000	
location_easting_osgr	101070.0	455370.590195	
location_northing_osgr	101070.0	274267.396596	
longitude	101070.0	-1.204951	
latitude	101070.0	52.355819	
police_force	101087.0	27.055833	
accident_severity	101087.0	2.760286	
number_of_vehicles	101087.0	1.844382	
number_of_casualties	101087.0	1.268304	
day_of_week	101087.0	4.138692	
local_authority_district	101087.0	-0.351756	
first_road_class	101087.0	4.204725	
first_road_number	101087.0	789.634899	
road_type	101087.0	5.256769	
speed_limit	101087.0	36.003146	
junction_detail	101087.0	4.385757	
junction_control	101087.0	1.771385	
second_road_class	101087.0	3.087697	
second_road_number	101087.0	223.660847	
pedestrian_crossing_human_control	101087.0	0.362064	
pedestrian_crossing_physical_facilities	101087.0	1.167648	

light_conditions	101087.0	1.976822
weather_conditions	101087.0	1.651419
road_surface_conditions	101087.0	1.346612
special_conditions_at_site	101087.0	0.253128
carriageway_hazards	101087.0	0.196623
urban_or_rural_area	101087.0	1.319695
did_police_officer_attend_scene_of_accident	101087.0	1.450533
trunk_road_flag	101087.0	1.721547

	std	min \
accident_year	0.000000	2021.000000
location_easting_osgr	92901.603757	67564.000000
location_northing_osgr	145053.875052	13898.000000
longitude	1.357913	-7.486852
latitude	1.307057	49.980835
police_force	24.235890	1.000000
accident_severity	0.459800	1.000000
number_of_vehicles	0.680355	1.000000
number_of_casualties	0.674544	1.000000
day_of_week	1.941061	1.000000
local_authority_district	15.603753	-1.000000
first_road_class	1.464199	1.000000
first_road_number	1583.067339	0.000000
road_type	1.701544	1.000000
speed_limit	14.121083	20.000000
junction_detail	13.915899	-1.000000
junction_control	2.528051	-1.000000
second_road_class	2.751656	-1.000000
second_road_number	931.658373	-1.000000
pedestrian_crossing_human_control	1.717190	-1.000000
pedestrian_crossing_physical_facilities	2.420433	-1.000000
light_conditions	1.685636	-1.000000
weather_conditions	1.848077	-1.000000
road_surface_conditions	0.957330	-1.000000
special_conditions_at_site	1.345088	-1.000000
carriageway_hazards	1.209476	-1.000000
urban_or_rural_area	0.466721	1.000000
did_police_officer_attend_scene_of_accident	0.729211	1.000000
trunk_road_flag	0.792752	-1.000000

	25%	50% \
accident_year	2021.000000	2021.000000
location_easting_osgr	392252.500000	461832.500000
location_northing_osgr	175175.250000	211641.500000
longitude	-2.116694	-1.087226
latitude	51.462397	51.789613
police_force	4.000000	21.000000

accident_severity	3.000000	3.000000
number_of_vehicles	1.000000	2.000000
number_of_casualties	1.000000	1.000000
day_of_week	2.000000	4.000000
local_authority_district	-1.000000	-1.000000
first_road_class	3.000000	4.000000
first_road_number	0.000000	34.000000
road_type	6.000000	6.000000
speed_limit	30.000000	30.000000
junction_detail	0.000000	2.000000
junction_control	-1.000000	2.000000
second_road_class	0.000000	3.000000
second_road_number	-1.000000	0.000000
pedestrian_crossing_human_control	0.000000	0.000000
pedestrian_crossing_physical_facilities	0.000000	0.000000
light_conditions	1.000000	1.000000
weather_conditions	1.000000	1.000000
road_surface_conditions	1.000000	1.000000
special_conditions_at_site	0.000000	0.000000
carriageway_hazards	0.000000	0.000000
urban_or_rural_area	1.000000	1.000000
did_police_officer_attend_scene_of_accident	1.000000	1.000000
trunk_road_flag	2.000000	2.000000

	75%	max
accident_year	2021.000000	2.021000e+03
location_easting_osgr	530038.000000	6.551400e+05
location_northing_osgr	382389.500000	1.179892e+06
longitude	-0.127483	1.755955e+00
latitude	53.334320	6.050001e+01
police_force	44.000000	9.900000e+01
accident_severity	3.000000	3.000000e+00
number_of_vehicles	2.000000	1.300000e+01
number_of_casualties	1.000000	2.200000e+01
day_of_week	6.000000	7.000000e+00
local_authority_district	-1.000000	4.800000e+02
first_road_class	6.000000	6.000000e+00
first_road_number	532.000000	9.480000e+03
road_type	6.000000	9.000000e+00
speed_limit	40.000000	7.000000e+01
junction_detail	3.000000	9.900000e+01
junction_control	4.000000	9.000000e+00
second_road_class	6.000000	6.000000e+00
second_road_number	0.000000	9.176000e+03
pedestrian_crossing_human_control	0.000000	9.000000e+00
pedestrian_crossing_physical_facilities	0.000000	9.000000e+00
light_conditions	4.000000	7.000000e+00

weather_conditions	1.000000	9.000000e+00
road_surface_conditions	2.000000	9.000000e+00
special_conditions_at_site	0.000000	9.000000e+00
carriageway_hazards	0.000000	9.000000e+00
urban_or_rural_area	2.000000	3.000000e+00
did_police_officer_attend_scene_of_accident	2.000000	3.000000e+00
trunk_road_flag	2.000000	2.000000e+00

```
[8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101087 entries, 0 to 101086
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   accident_index                        101087 non-null object
1   accident_year                         101087 non-null int64
2   accident_reference                    101087 non-null object
3   location_easting_osgr                 101070 non-null float64
4   location_northing_osgr                101070 non-null float64
5   longitude                             101070 non-null float64
6   latitude                              101070 non-null float64
7   police_force                          101087 non-null int64
8   accident_severity                     101087 non-null int64
9   number_of_vehicles                    101087 non-null int64
10  number_of_casualties                  101087 non-null int64
11  date                                  101087 non-null object
12  day_of_week                           101087 non-null int64
13  time                                  101087 non-null object
14  local_authority_district               101087 non-null int64
15  local_authority_ons_district           101087 non-null object
16  local_authority_highway                101087 non-null object
17  first_road_class                       101087 non-null int64
18  first_road_number                      101087 non-null int64
19  road_type                              101087 non-null int64
20  speed_limit                            101087 non-null int64
21  junction_detail                       101087 non-null int64
22  junction_control                       101087 non-null int64
23  second_road_class                     101087 non-null int64
24  second_road_number                    101087 non-null int64
25  pedestrian_crossing_human_control      101087 non-null int64
26  pedestrian_crossing_physical_facilities 101087 non-null int64
27  light_conditions                       101087 non-null int64
28  weather_conditions                    101087 non-null int64
29  road_surface_conditions                101087 non-null int64
30  special_conditions_at_site             101087 non-null int64
31  carriageway_hazards                   101087 non-null int64
```



```

32  urban_or_rural_area          101087 non-null  int64
33  did_police_officer_attend_scene_of_accident  101087 non-null  int64
34  trunk_road_flag             101087 non-null  int64
35  lsoa_of_accident_location    101087 non-null  object
dtypes: float64(4), int64(25), object(7)
memory usage: 27.8+ MB

```

```

[9]: # Checking for missing values in the data
data.isnull().sum()

```

```

[9]: accident_index          0
accident_year              0
accident_reference         0
location_easting_osgr      17
location_northing_osgr     17
longitude                 17
latitude                 17
police_force              0
accident_severity         0
number_of_vehicles        0
number_of_casualties      0
date                     0
day_of_week              0
time                     0
local_authority_district  0
local_authority_ons_district 0
local_authority_highway   0
first_road_class          0
first_road_number         0
road_type                 0
speed_limit               0
junction_detail           0
junction_control          0
second_road_class         0
second_road_number        0
pedestrian_crossing_human_control 0
pedestrian_crossing_physical_facilities 0
light_conditions          0
weather_conditions        0
road_surface_conditions   0
special_conditions_at_site 0
carriageway_hazards       0
urban_or_rural_area       0
did_police_officer_attend_scene_of_accident 0
trunk_road_flag           0
lsoa_of_accident_location 0
dtype: int64

```

```
[31]: data.nunique()
```

```
[31]: police_force          44
      accident_severity     3
      number_of_vehicles    13
      number_of_casualties  12
      day_of_week           7
      local_authority_district 15
      first_road_class       6
      first_road_number    3099
      road_type             6
      speed_limit           6
      junction_detail       11
      junction_control      6
      second_road_class     8
      second_road_number   2374
      pedestrian_crossing_human_control 5
      pedestrian_crossing_physical_facilities 8
      light_conditions      6
      weather_conditions    10
      road_surface_conditions 7
      special_conditions_at_site 10
      carriageway_hazards   8
      urban_or_rural_area   3
      did_police_officer_attend_scene_of_accident 3
      trunk_road_flag       3
      dtype: int64
```

```
[12]: # deleting columns which are mostly unique or has a unique category
      data = data.
      ↪drop(['accident_index', 'accident_year', 'accident_reference', 'location_easting_osgr',
            ↵
            ↪'location_northing_osgr', 'longitude', 'latitude', 'date', 'time', 'lsoa_of_accident_location',
            ↵
            ↪'local_authority_ons_district', 'local_authority_highway'], ↵
      ↪axis = 1)
```

```
[17]: data.duplicated().sum()
```

```
[17]: 0
```

```
[16]: # dropping duplicates
      data = data.drop_duplicates()
```

```
[21]: # converting the datatypes of categorical variables to category
      data['police_force'] = data['police_force'].astype('category')
      data['day_of_week'] = data['day_of_week'].astype('category')
      data['accident_severity'] = data['accident_severity'].astype('category')
```

```

data['first_road_class'] = data['first_road_class'].astype('category')
data['local_authority_district'] = data['local_authority_district'].
    ↳astype('category')
data['second_road_class'] = data['second_road_class'].astype('category')

data['pedestrian_crossing_human_control'] =
    ↳data['pedestrian_crossing_human_control'].astype('category')
data['pedestrian_crossing_physical_facilities'] =
    ↳data['pedestrian_crossing_physical_facilities'].astype('category')
data['light_conditions'] = data['light_conditions'].astype('category')
data['weather_conditions'] = data['weather_conditions'].astype('category')
data['special_conditions_at_site'] = data['special_conditions_at_site'].
    ↳astype('category')
data['carriageway_hazards'] = data['carriageway_hazards'].astype('category')
data['did_police_officer_attend_scene_of_accident'] =
    ↳data['did_police_officer_attend_scene_of_accident'].astype('category')

data['trunk_road_flag'] = data['trunk_road_flag'].astype('category')
data['road_surface_conditions'] = data['road_surface_conditions'].
    ↳astype('category')
data['urban_or_rural_area'] = data['urban_or_rural_area'].astype('category')

```

```

[22]: num_cols =
    ↳['number_of_vehicles', 'number_of_casualties', 'first_road_number', 'speed_limit',
    ↳'junction_detail', 'junction_control']
cat_cols =
    ↳['police_force', 'day_of_week', 'accident_severity', 'first_road_class', 'local_authority_district',
    ↳'second_road_class', 'pedestrian_crossing_human_control', 'pedestrian_crossing_physical_facilities',
    ↳'weather_conditions', 'special_conditions_at_site', 'carriageway_hazards', 'did_police_officer_attend_scene_of_accident',
    ↳'road_surface_conditions', 'trunk_road_flag', 'urban_or_rural_area']

```

```

[24]: # let's check for missing values in the data
round(data.isnull().sum() / data.isnull().count() * 100, 2)

```

```

[24]: police_force          0.0
      accident_severity    0.0
      number_of_vehicles    0.0
      number_of_casualties  0.0
      day_of_week          0.0
      local_authority_district 0.0
      first_road_class      0.0
      first_road_number     0.0
      road_type             0.0
      speed_limit           0.0

```

junction_detail	0.0
junction_control	0.0
second_road_class	0.0
second_road_number	0.0
pedestrian_crossing_human_control	0.0
pedestrian_crossing_physical_facilities	0.0
light_conditions	0.0
weather_conditions	0.0
road_surface_conditions	0.0
special_conditions_at_site	0.0
carriageway_hazards	0.0
urban_or_rural_area	0.0
did_police_officer_attend_scene_of_accident	0.0
trunk_road_flag	0.0
dtype: float64	

```
[25]: # checking for value_counts in each nominal feature
columns = cat_cols
for i in columns:
    print("Unique values in", i, "are :")
    print(data[i].value_counts())
    print("*" * 50)
```

Unique values in police\_force are :

1	21738
20	4094
99	3749
13	3532
46	3423
47	3006
6	2815
44	2808
43	2770
42	2601
4	2592
50	2570
45	2387
52	2168
14	1942
5	1804
31	1750
30	1706
10	1706
16	1683
7	1546
22	1525
41	1517

```

32      1490
36      1315
35      1282
12      1251
55      1232
54      1223
33      1086
40      1038
37      1017
34      1013
63       947
62       902
23       832
21       806
60       766
53       737
3        724
11       623
17       602
61       590
48       145
Name: police_force, dtype: int64
*****
Unique values in day_of_week are :
6      15945
5      14149
4      13974
3      13680
7      13256
2      13117
1      10932
Name: day_of_week, dtype: int64
*****
Unique values in accident_severity are :
3      73048
2      20533
1      1472
Name: accident_severity, dtype: int64
*****
Unique values in first_road_class are :
3      44706
6      30379
4      12580
5       4123
1       2988
2        277
Name: first_road_class, dtype: int64
*****

```

Unique values in local\_authority\_district are :

-1	94859
241	32
245	31
243	27
240	25
472	14
476	14
471	10
480	9
470	7
474	7
473	6
477	5
475	4
478	3

Name: local\_authority\_district, dtype: int64

\*\*\*\*\*

Unique values in second\_road\_class are :

6	38843
0	38307
3	10774
4	3918
5	2767
1	337
-1	84
2	23

Name: second\_road\_class, dtype: int64

\*\*\*\*\*

Unique values in pedestrian\_crossing\_human\_control are :

0	89432
9	3720
2	1217
1	437
-1	247

Name: pedestrian\_crossing\_human\_control, dtype: int64

\*\*\*\*\*

Unique values in pedestrian\_crossing\_physical\_facilities are :

0	70749
5	8022
4	5679
1	4459
9	3140
8	2508
7	264
-1	232

Name: pedestrian\_crossing\_physical\_facilities, dtype: int64

\*\*\*\*\*

```

Unique values in light_conditions are :
1      68118
4      19401
6       4840
7       1972
5        711
-1        11
Name: light_conditions, dtype: int64
*****
Unique values in weather_conditions are :
1      75824
2      10290
8       3203
9       2698
5       1033
4        872
3        603
7        427
6         91
-1        12
Name: weather_conditions, dtype: int64
*****
Unique values in special_conditions_at_site are :
0      90447
9       1742
4       1406
-1       351
1        318
5        248
7        203
3        157
6        130
2         51
Name: special_conditions_at_site, dtype: int64
*****
Unique values in carriageway_hazards are :
0      91075
9       1455
2       1184
7        348
-1       342
1        262
6        202
3        185
Name: carriageway_hazards, dtype: int64
*****
Unique values in did_police_officer_attend_scene_of_accident are :
1      65704

```

```

2    15612
3    13737
Name: did_police_officer_attend_scene_of_accident, dtype: int64
*****
Unique values in road_surface_conditions are :
1    69488
2    22327
4     1202
9     1001
3      473
-1     428
5      134
Name: road_surface_conditions, dtype: int64
*****
Unique values in trunk_road_flag are :
2    81736
-1    6954
1     6363
Name: trunk_road_flag, dtype: int64
*****
Unique values in urban_or_rural_area are :
1    63633
2    31403
3         17
Name: urban_or_rural_area, dtype: int64
*****

```

```

[26]: # Checking summary statistics
data[num_cols].describe().T

```

```

[26]:

```

	count	mean	std	min	25%	50%	\
number_of_vehicles	95053.0	1.850462	0.690398	1.0	1.0	2.0	
number_of_casualties	95053.0	1.282737	0.690961	1.0	1.0	1.0	
first_road_number	95053.0	836.641505	1617.922890	0.0	0.0	46.0	
speed_limit	95053.0	36.287334	14.262362	20.0	30.0	30.0	
junction_detail	95053.0	4.464825	14.038640	-1.0	0.0	2.0	
junction_control	95053.0	1.794167	2.519444	-1.0	-1.0	2.0	

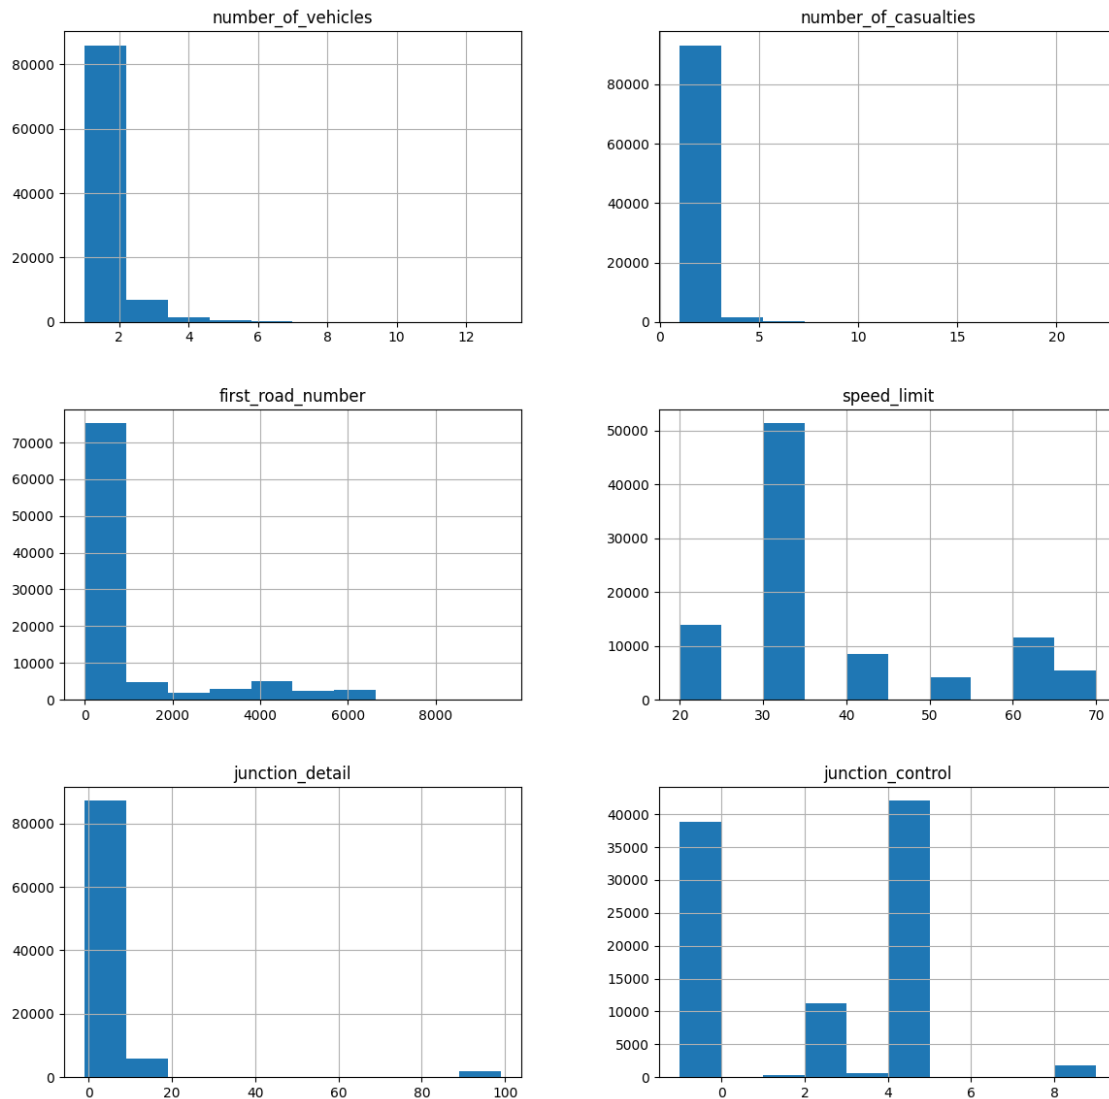
  

	75%	max
number_of_vehicles	2.0	13.0
number_of_casualties	1.0	22.0
first_road_number	595.0	9480.0
speed_limit	40.0	70.0
junction_detail	3.0	99.0
junction_control	4.0	9.0



```
[28]: # Creating histograms
data[num_cols].hist(figsize = (14, 14))

plt.show()
```



```
[ ]:
```

```
[ ]:
```

```
[30]: data.head()
```

```
[30]:  police_force  accident_severity  number_of_vehicles  number_of_casualties  \
0           0           1           3           3           1
```

1	1	2	2	3
2	1	2	2	4
3	1	1	1	1
4	1	3	4	1

	day_of_week	local_authority_district	first_road_class	first_road_number	\
0	6	-1	6	0	
1	6	-1	3	1203	
2	6	-1	4	272	
3	6	-1	3	3220	
4	6	-1	5	0	

	road_type	speed_limit	junction_detail	junction_control	\
0	6	30	9	4	
1	3	30	7	2	
2	6	30	9	2	
3	2	30	9	4	
4	6	20	3	4	

	second_road_class	second_road_number	pedestrian_crossing_human_control	\
0	6	0	0	
1	3	1204	0	
2	5	0	0	
3	6	0	0	
4	6	0	0	

	pedestrian_crossing_physical_facilities	light_conditions	weather_conditions	\
0		0	4	7
1		5	4	1
2		5	4	1
3		4	4	1
4		0	4	1

	road_surface_conditions	special_conditions_at_site	carriageway_hazards	\
0	4		1	0
1	1		0	0
2	1		0	0
3	1		0	0
4	1		0	0

	urban_or_rural_area	did_police_officer_attend_scene_of_accident	\
0	1	1	
1	1	1	
2	1	1	
3	1	1	
4	1	1	

```

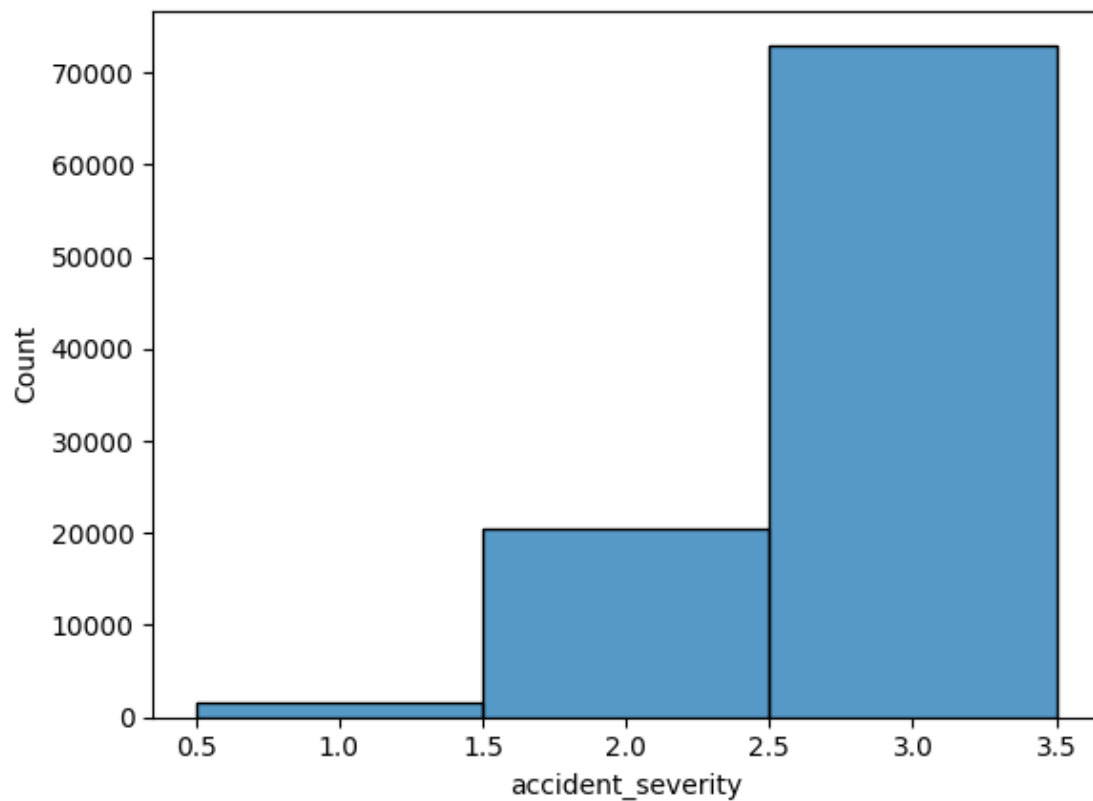
trunk_road_flag
0          2
1          2
2          2
3          2
4          2

```

```

[33]: # Distribution of target variable
sns.histplot(data["accident_severity"], kde=False)
plt.show()

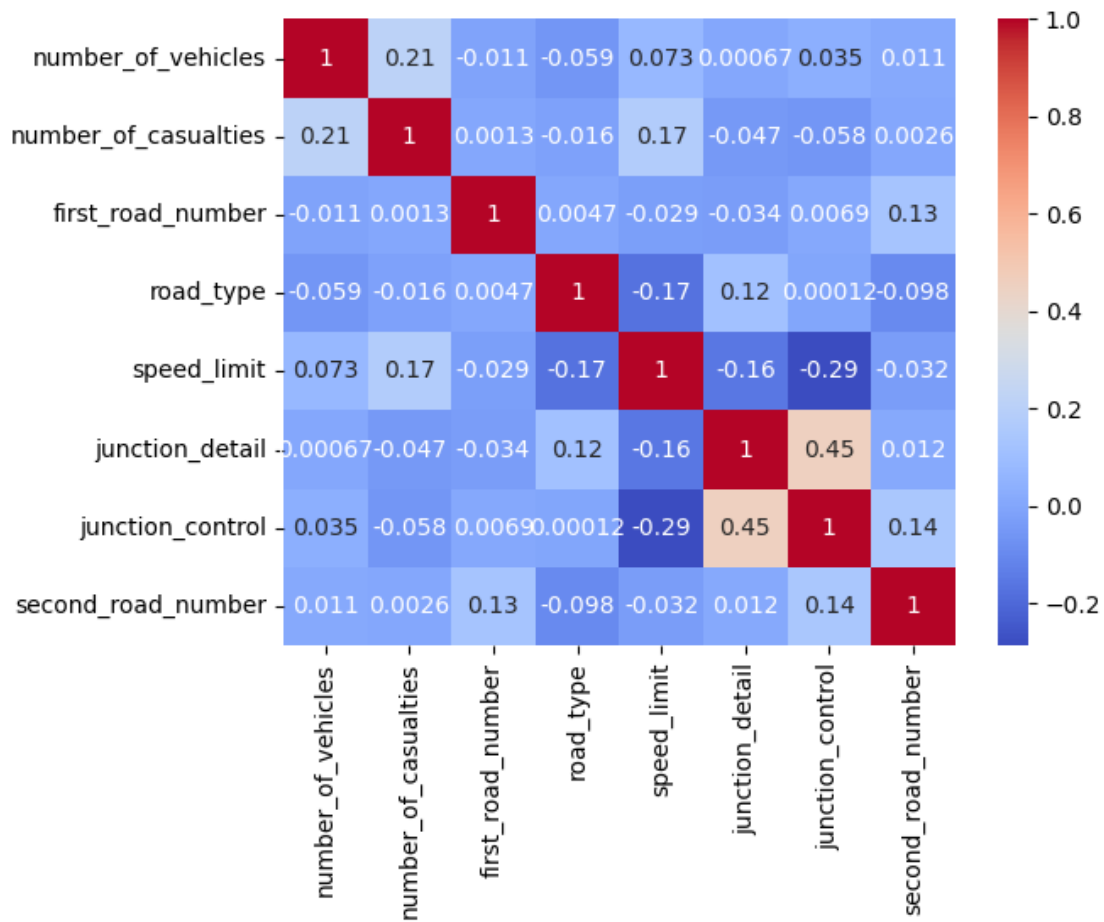
```



```

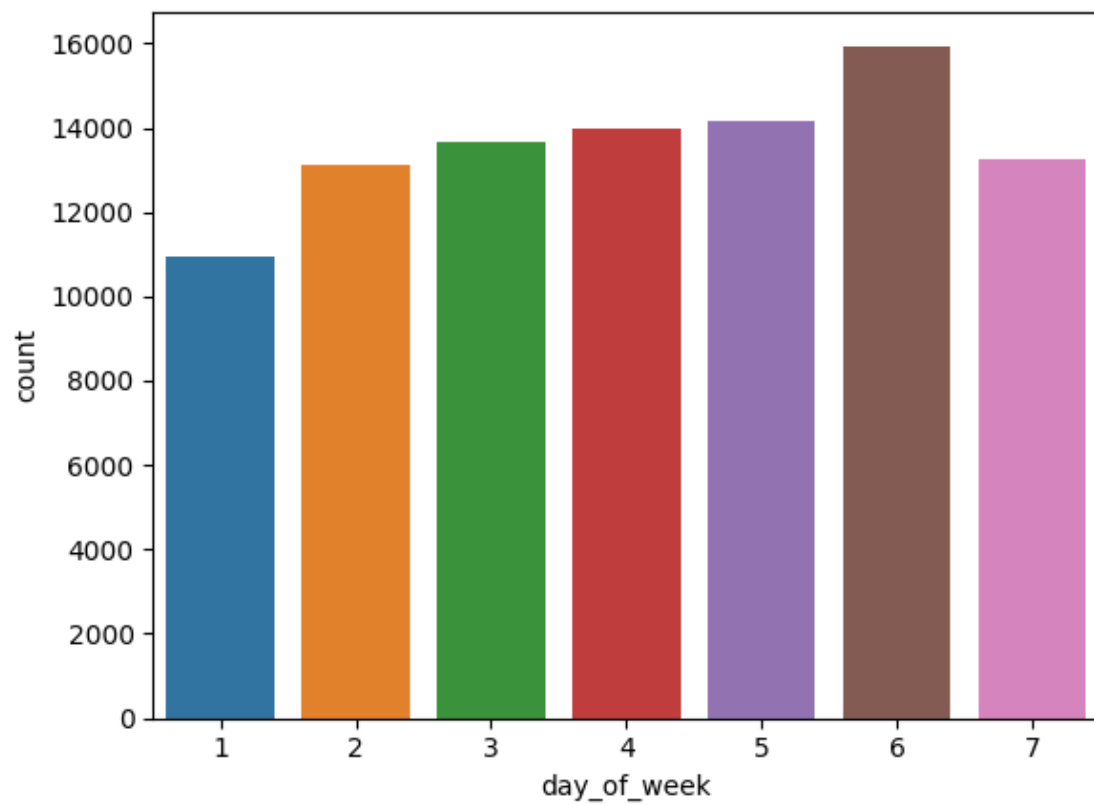
[35]: # Correlation analysis
corr = data.corr()
sns.heatmap(corr, cmap="coolwarm", annot=True)
plt.show()

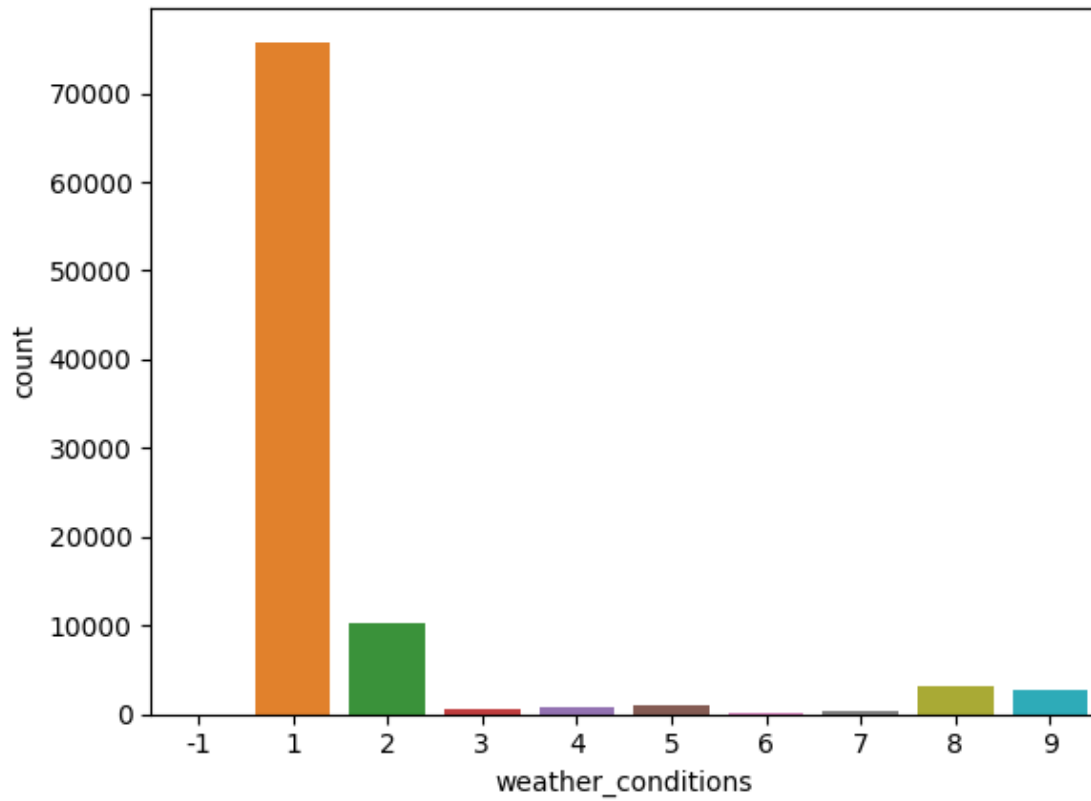
```



```
[36]: # Bar chart for categorical variables
sns.countplot(x="day_of_week", data=data)
plt.show()

sns.countplot(x="weather_conditions", data=data)
plt.show()
```





```
[38]: for i in cat_cols:
      print(data[i].value_counts(normalize = True))

      print('*' * 40)
```

```
1      0.228693
20     0.043071
99     0.039441
13     0.037158
46     0.036011
47     0.031624
6      0.029615
44     0.029541
43     0.029142
42     0.027364
4      0.027269
50     0.027038
45     0.025112
52     0.022808
14     0.020431
5      0.018979
```

```

31    0.018411
30    0.017948
10    0.017948
16    0.017706
7     0.016265
22    0.016044
41    0.015960
32    0.015675
36    0.013834
35    0.013487
12    0.013161
55    0.012961
54    0.012867
33    0.011425
40    0.010920
37    0.010699
34    0.010657
63    0.009963
62    0.009489
23    0.008753
21    0.008479
60    0.008059
53    0.007754
3     0.007617
11    0.006554
17    0.006333
61    0.006207
48    0.001525
Name: police_force, dtype: float64
*****
6     0.167749
5     0.148854
4     0.147013
3     0.143920
7     0.139459
2     0.137997
1     0.115010
Name: day_of_week, dtype: float64
*****
3     0.768498
2     0.216016
1     0.015486
Name: accident_severity, dtype: float64
*****
3     0.470327
6     0.319601
4     0.132347
5     0.043376

```

```

1      0.031435
2      0.002914
Name: first_road_class, dtype: float64
*****
-1      0.997959
241     0.000337
245     0.000326
243     0.000284
240     0.000263
472     0.000147
476     0.000147
471     0.000105
480     0.000095
470     0.000074
474     0.000074
473     0.000063
477     0.000053
475     0.000042
478     0.000032
Name: local_authority_district, dtype: float64
*****
6       0.408646
0       0.403007
3       0.113347
4       0.041219
5       0.029110
1       0.003545
-1      0.000884
2       0.000242
Name: second_road_class, dtype: float64
*****
0       0.940865
9       0.039136
2       0.012803
1       0.004597
-1      0.002599
Name: pedestrian_crossing_human_control, dtype: float64
*****
0       0.744311
5       0.084395
4       0.059746
1       0.046911
9       0.033034
8       0.026385
7       0.002777
-1      0.002441
Name: pedestrian_crossing_physical_facilities, dtype: float64
*****

```



```

1      0.716632
4      0.204107
6      0.050919
7      0.020746
5      0.007480
-1     0.000116
Name: light_conditions, dtype: float64
*****
1      0.797702
2      0.108255
8      0.033697
9      0.028384
5      0.010868
4      0.009174
3      0.006344
7      0.004492
6      0.000957
-1     0.000126
Name: weather_conditions, dtype: float64
*****
0      0.951543
9      0.018327
4      0.014792
-1     0.003693
1      0.003346
5      0.002609
7      0.002136
3      0.001652
6      0.001368
2      0.000537
Name: special_conditions_at_site, dtype: float64
*****
0      0.958150
9      0.015307
2      0.012456
7      0.003661
-1     0.003598
1      0.002756
6      0.002125
3      0.001946
Name: carriageway_hazards, dtype: float64
*****
1      0.691235
2      0.164245
3      0.144519
Name: did_police_officer_attend_scene_of_accident, dtype: float64
*****
1      0.731045

```

```

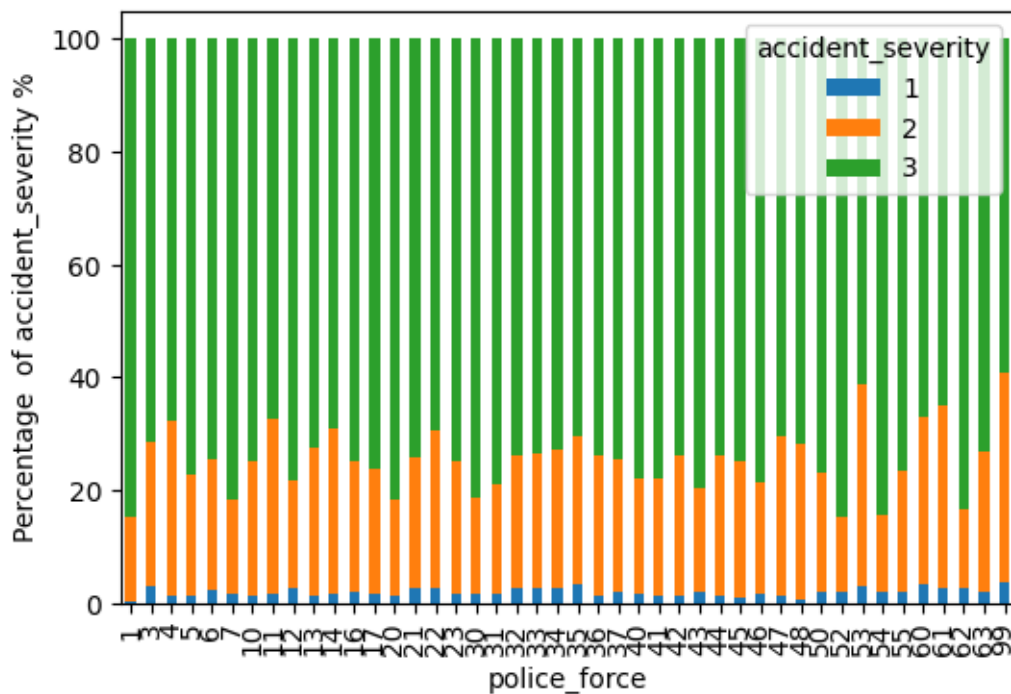
2      0.234890
4      0.012646
9      0.010531
3      0.004976
-1     0.004503
5      0.001410
Name: road_surface_conditions, dtype: float64
*****
2      0.859899
-1     0.073159
1      0.066942
Name: trunk_road_flag, dtype: float64
*****
1      0.669448
2      0.330374
3      0.000179
Name: urban_or_rural_area, dtype: float64
*****

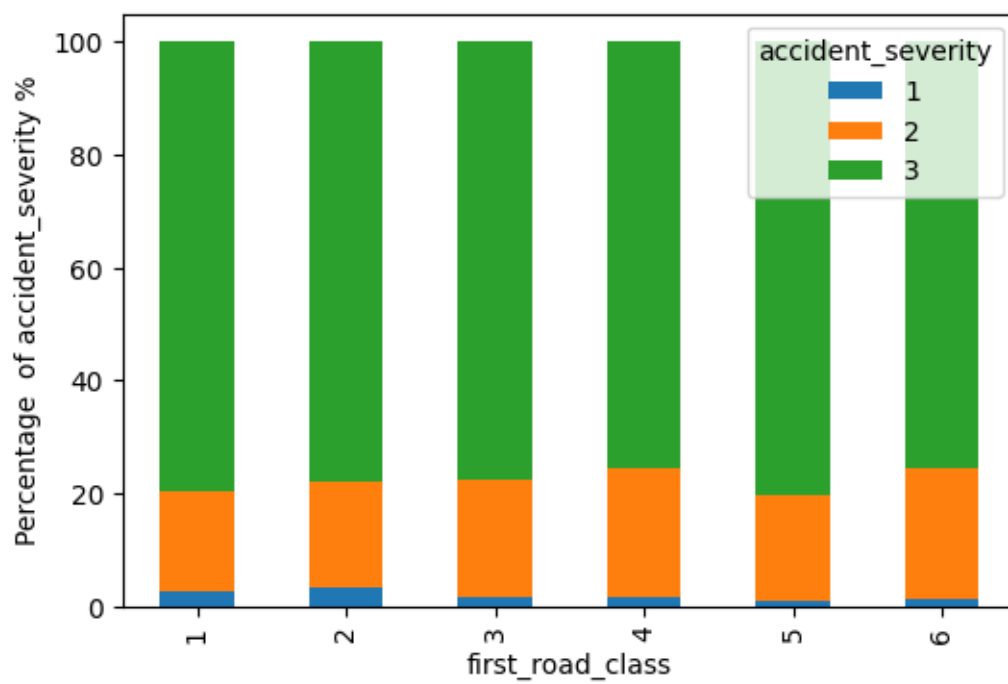
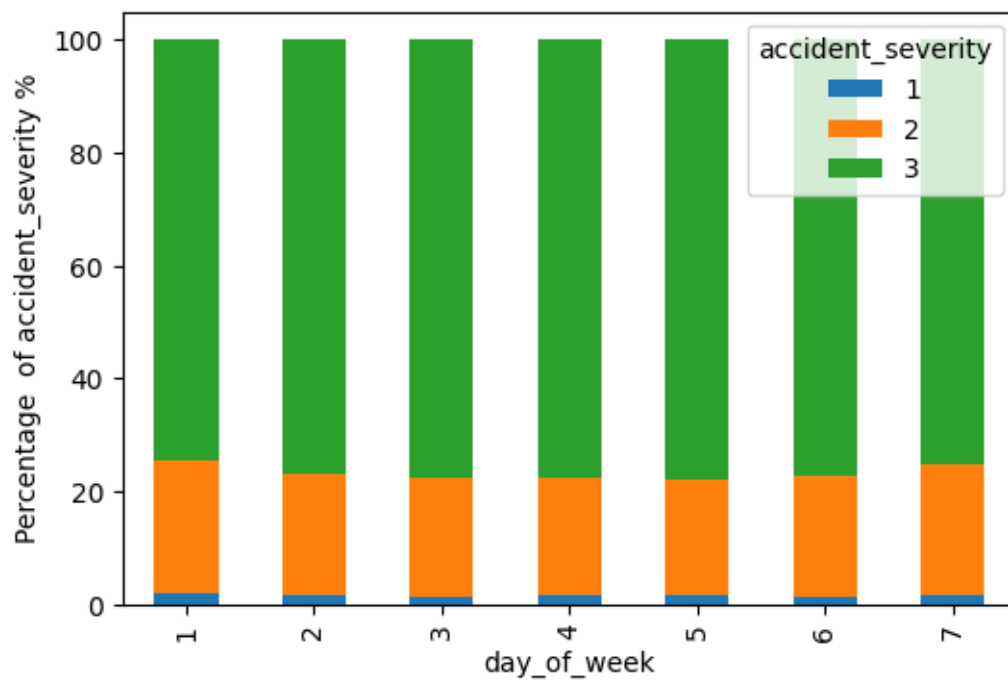
```

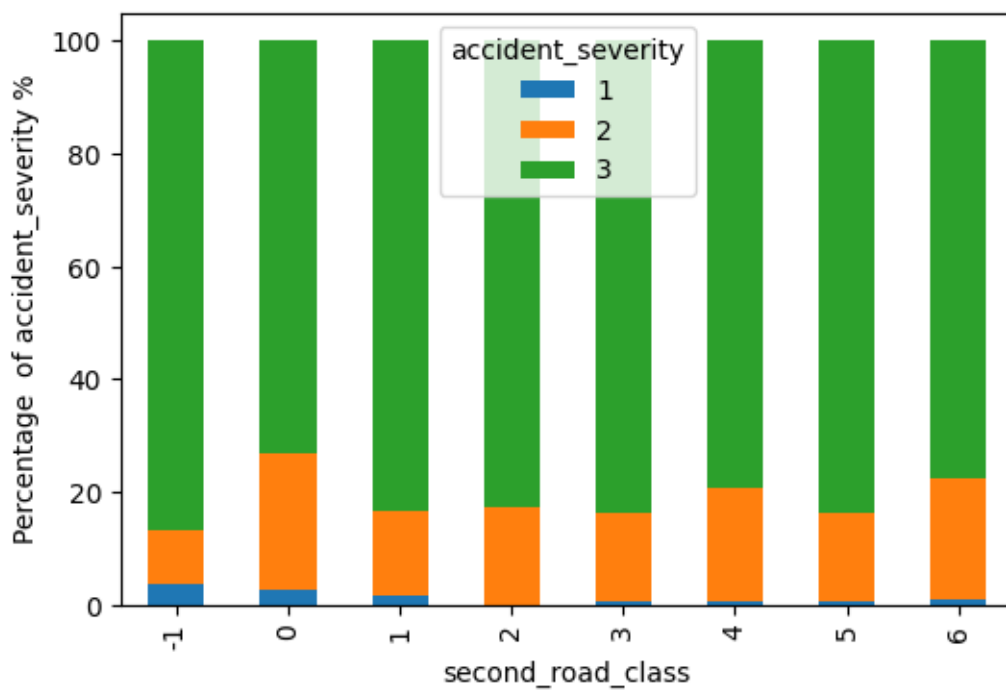
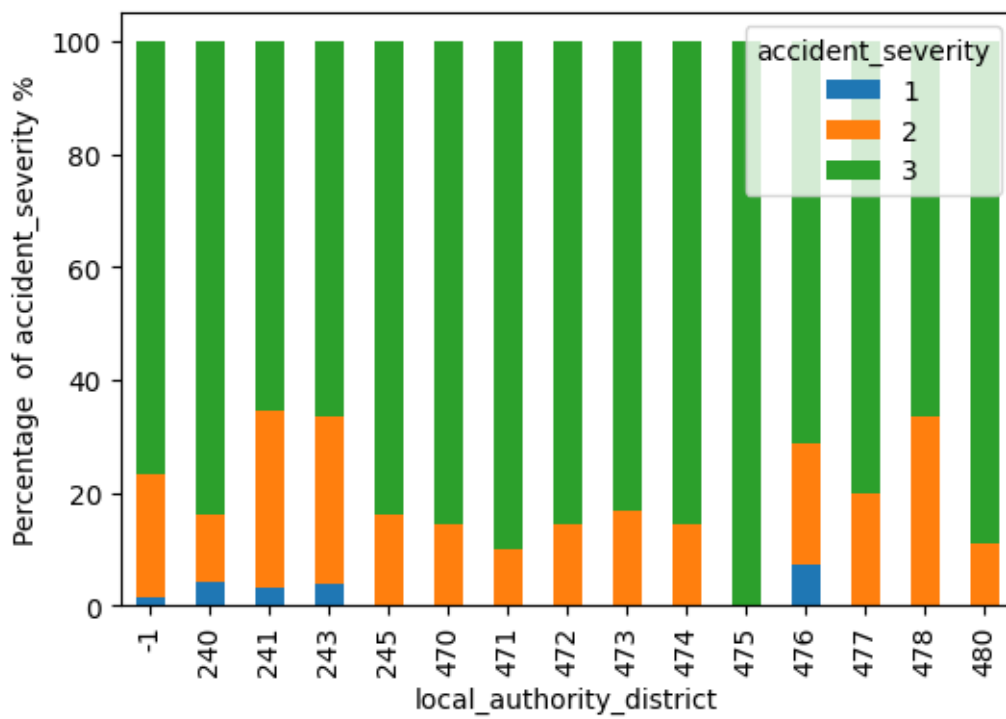
```

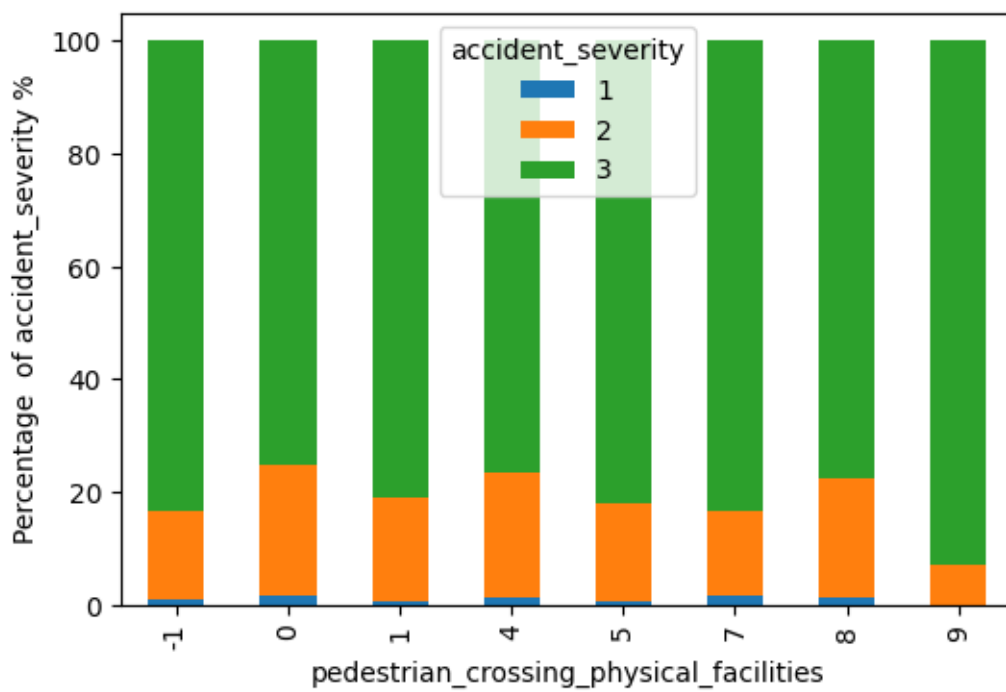
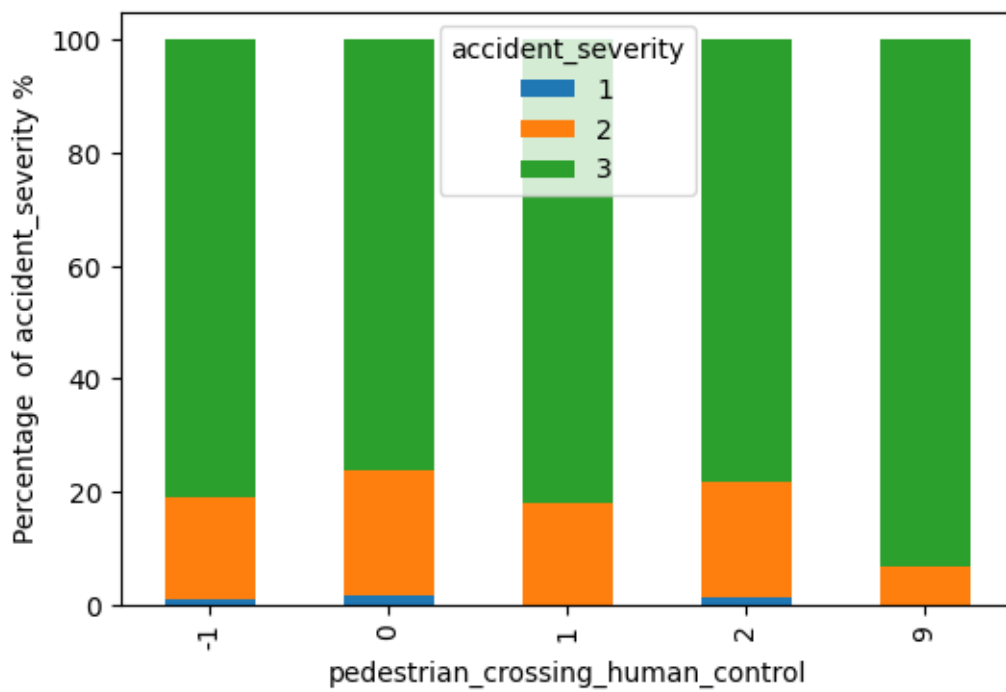
[39]: for i in cat_cols:
        if i != 'accident_severity':
            (pd.crosstab(data[i], data['accident_severity'], normalize =
↳ 'index')*100).plot(kind = 'bar', figsize = (6, 4), stacked = True)
            plt.ylabel('Percentage of accident_severity %')

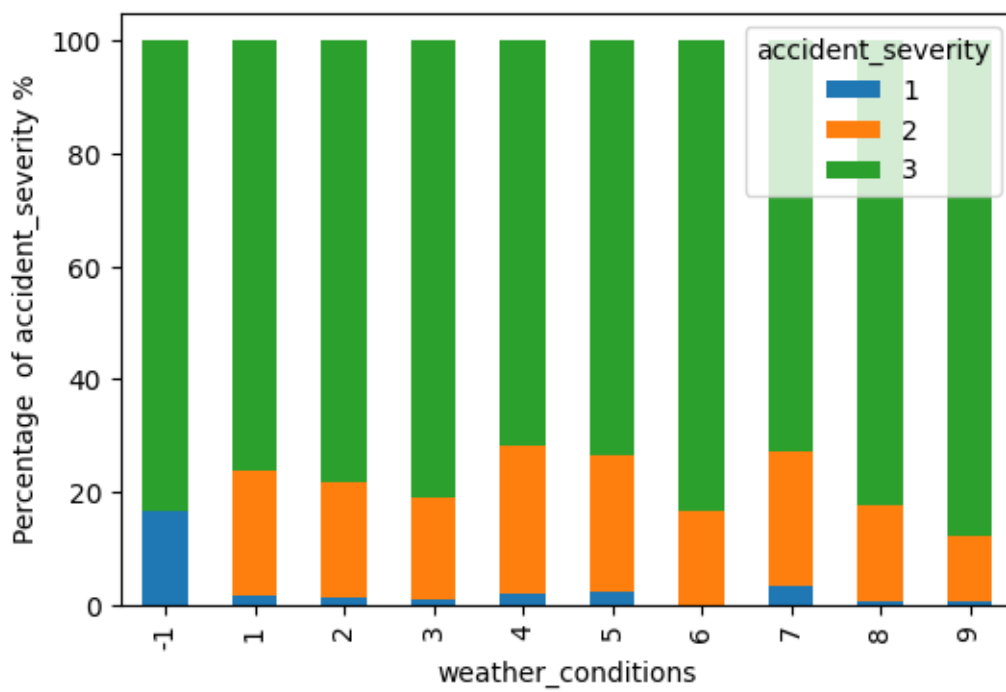
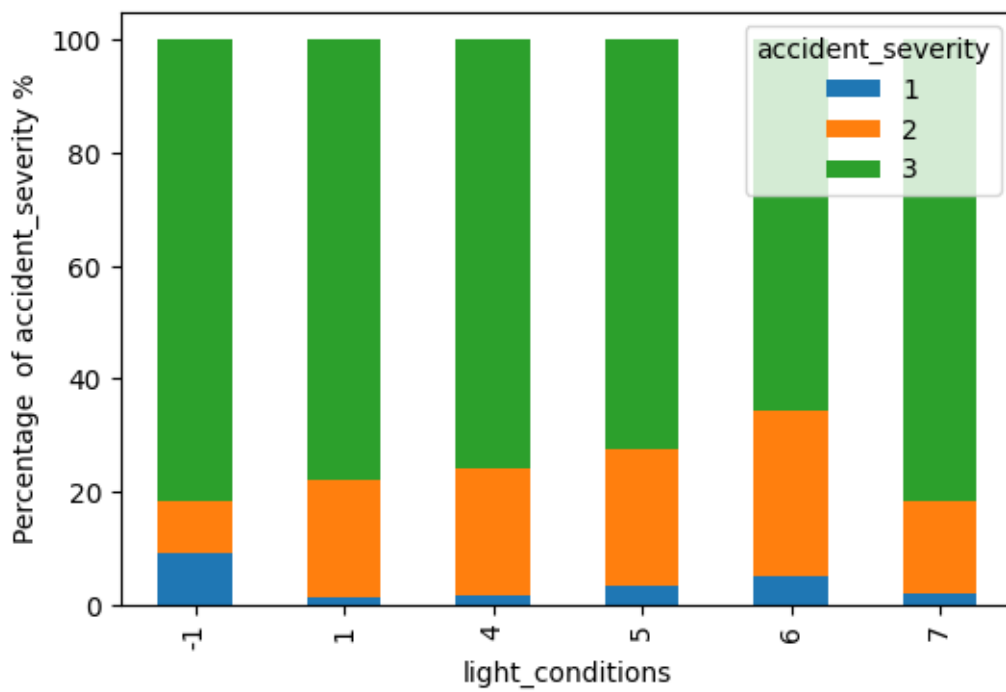
```

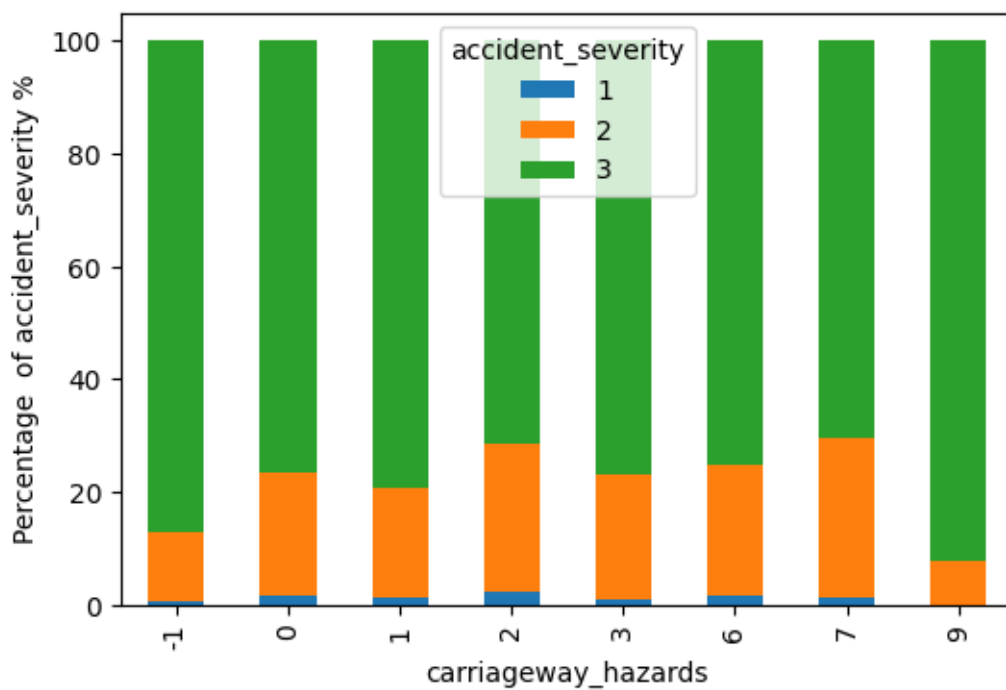
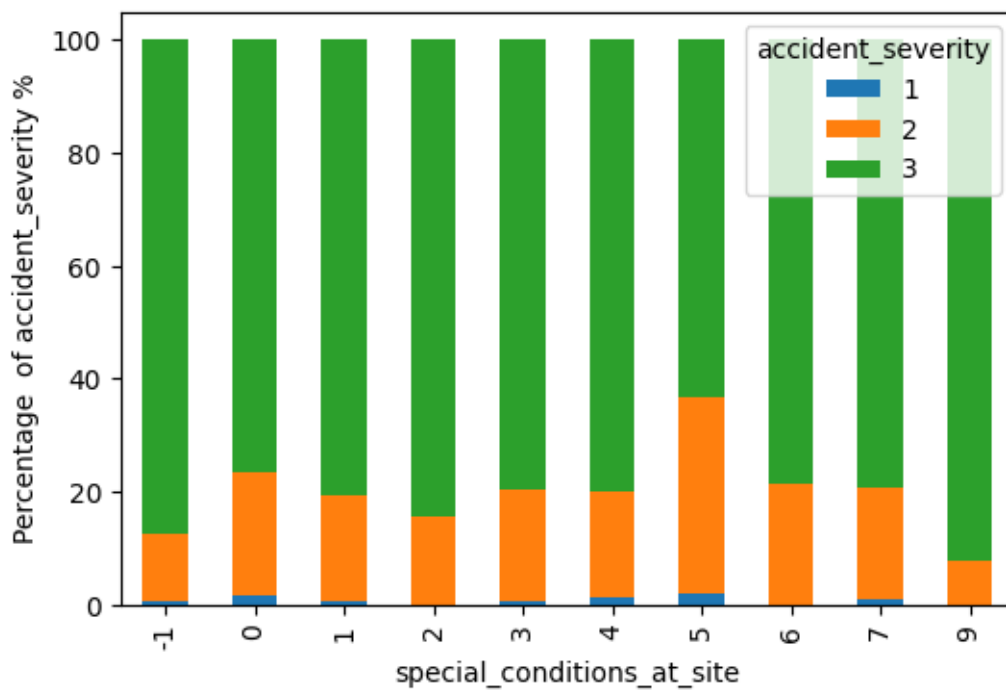


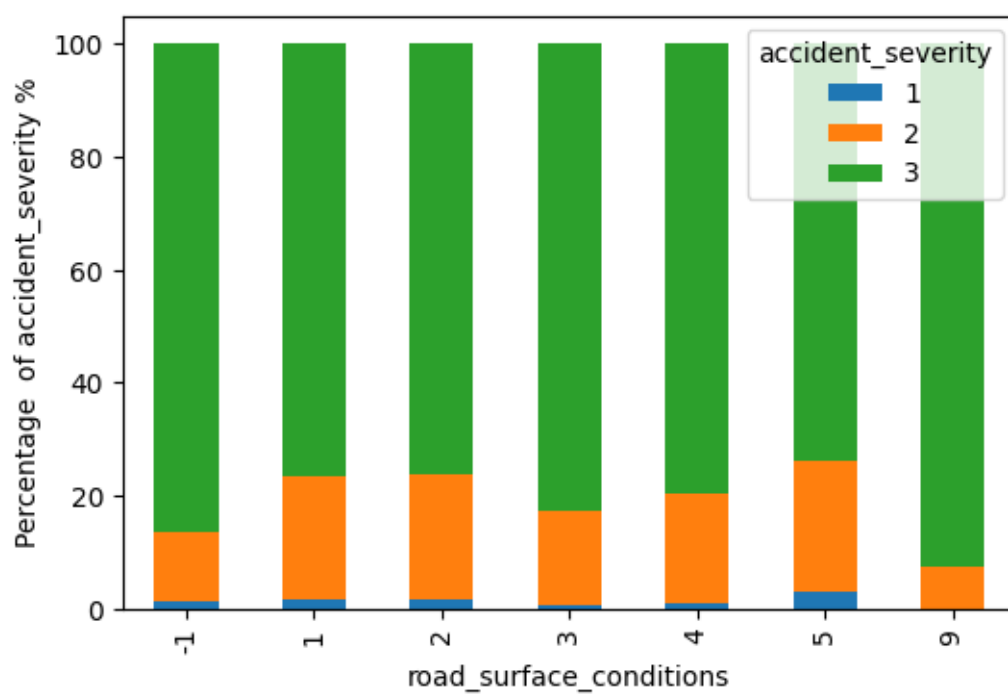
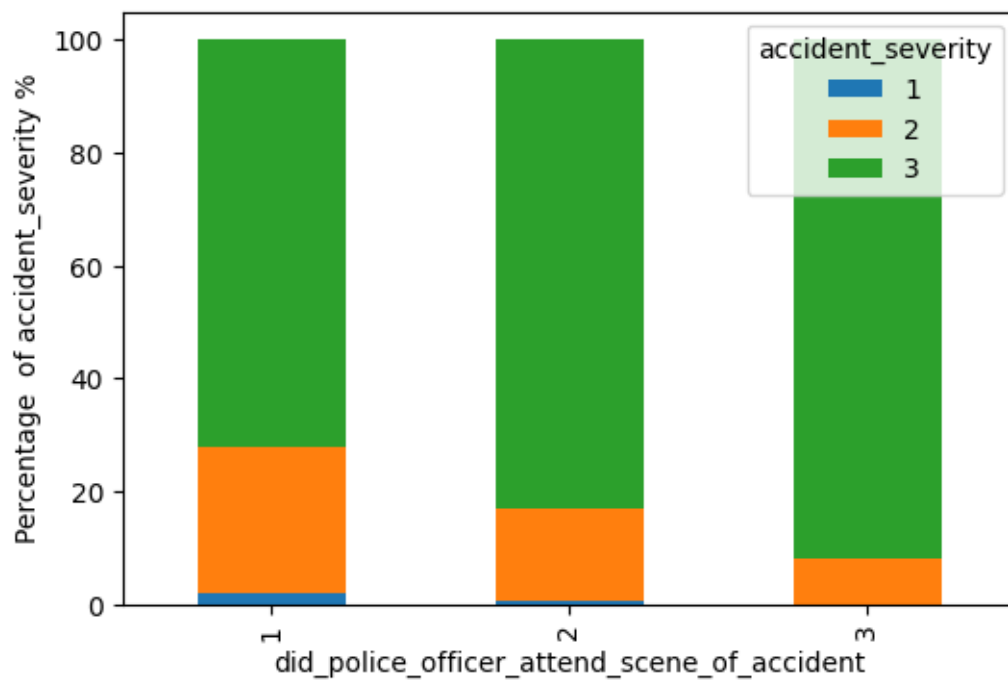




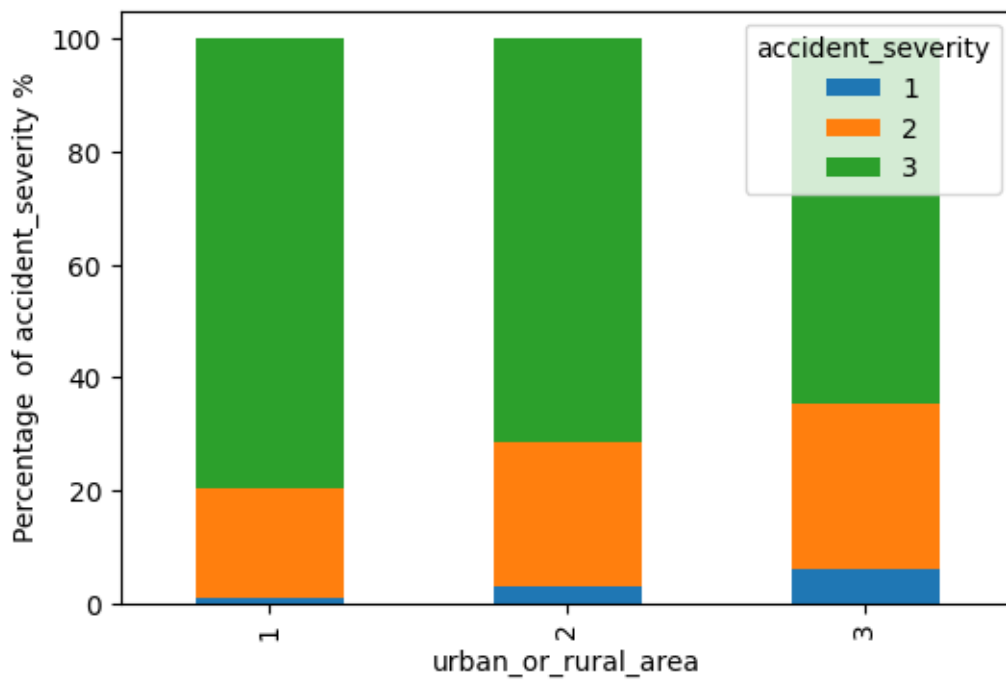
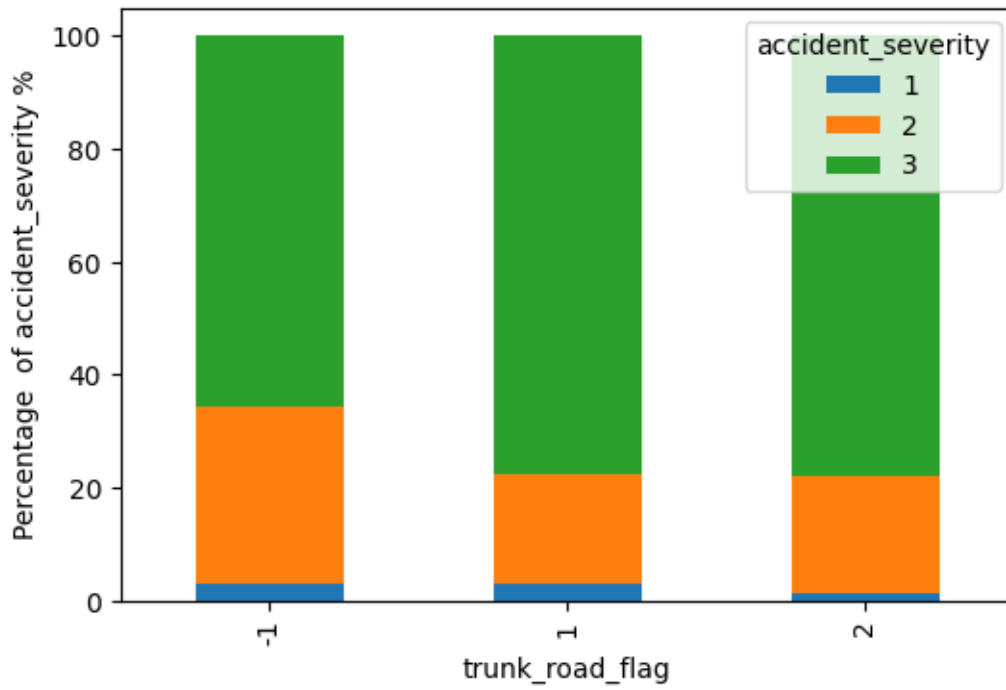












[ ]:

```
[41]: data[num_cols].corr()
```

```
[41]:
```

	number_of_vehicles	number_of_casualties	\
number_of_vehicles	1.000000	0.212728	
number_of_casualties	0.212728	1.000000	
first_road_number	-0.010713	0.001285	
speed_limit	0.073207	0.173840	
junction_detail	0.000665	-0.046740	
junction_control	0.034882	-0.057680	

	first_road_number	speed_limit	junction_detail	\
number_of_vehicles	-0.010713	0.073207	0.000665	
number_of_casualties	0.001285	0.173840	-0.046740	
first_road_number	1.000000	-0.029029	-0.033807	
speed_limit	-0.029029	1.000000	-0.155483	
junction_detail	-0.033807	-0.155483	1.000000	
junction_control	0.006942	-0.285929	0.454381	

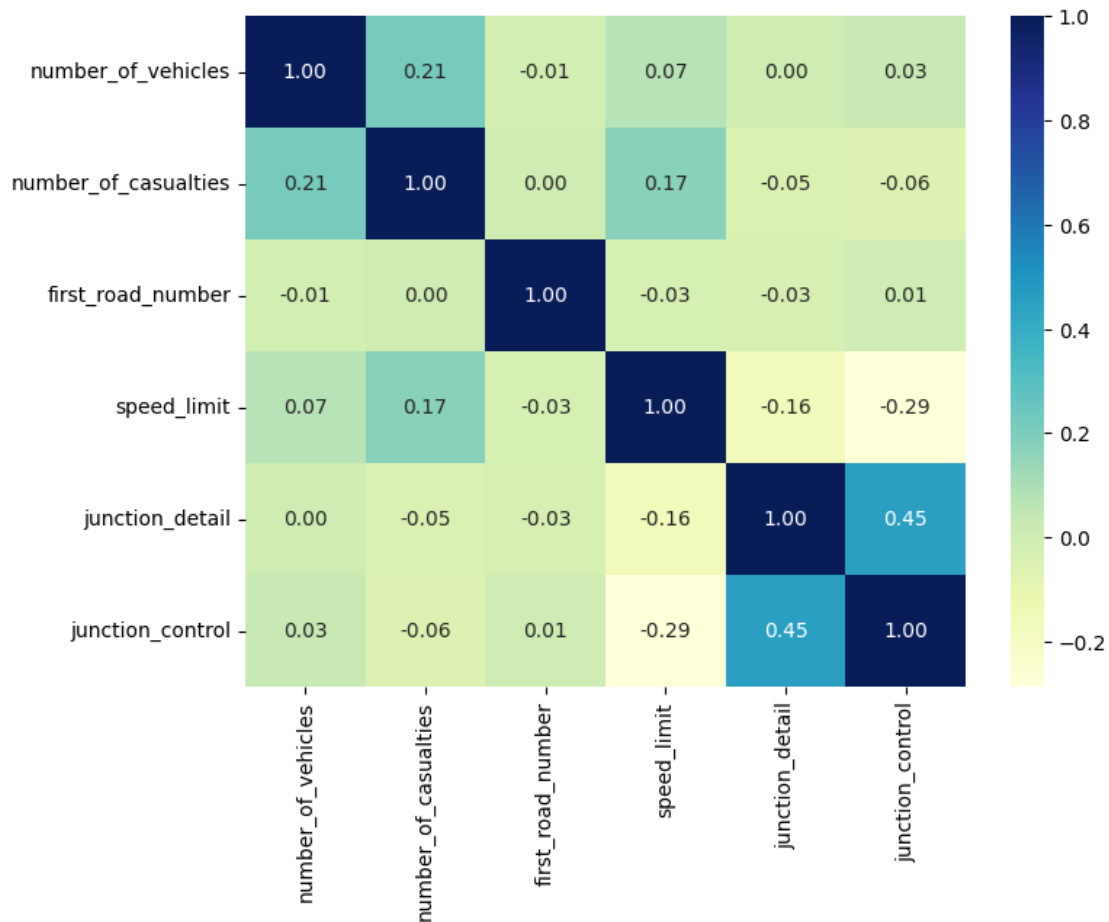
  

	junction_control
number_of_vehicles	0.034882
number_of_casualties	-0.057680
first_road_number	0.006942
speed_limit	-0.285929
junction_detail	0.454381
junction_control	1.000000

```
[40]: # Plotting the correlation between numerical variables
plt.figure(figsize = (8, 6))

sns.heatmap(data[num_cols].corr(), annot = True, fmt = '0.2f', cmap = 'YlGnBu')
```

```
[40]: <Axes: >
```



```
[43]: # Function to plot a boxplot and a histogram along the same scale

def histogram_boxplot(data, feature, figsize = (10, 7), kde = False, bins =
↳None):
    f2, (ax_box2, ax_hist2) = plt.subplots(
        nrows = 2,
        sharex = True,
        gridspec_kw = {"height_ratios": (0.25, 0.75)},
        figsize = figsize,
    )
    sns.boxplot(
        data = data, x = feature, ax = ax_box2, showmeans = True, color =
↳"violet"
    )
    sns.histplot(
        data = data, x = feature, kde = kde, ax = ax_hist2, bins = bins,
↳palette = "winter"
    ) if bins else sns.histplot(
```

```

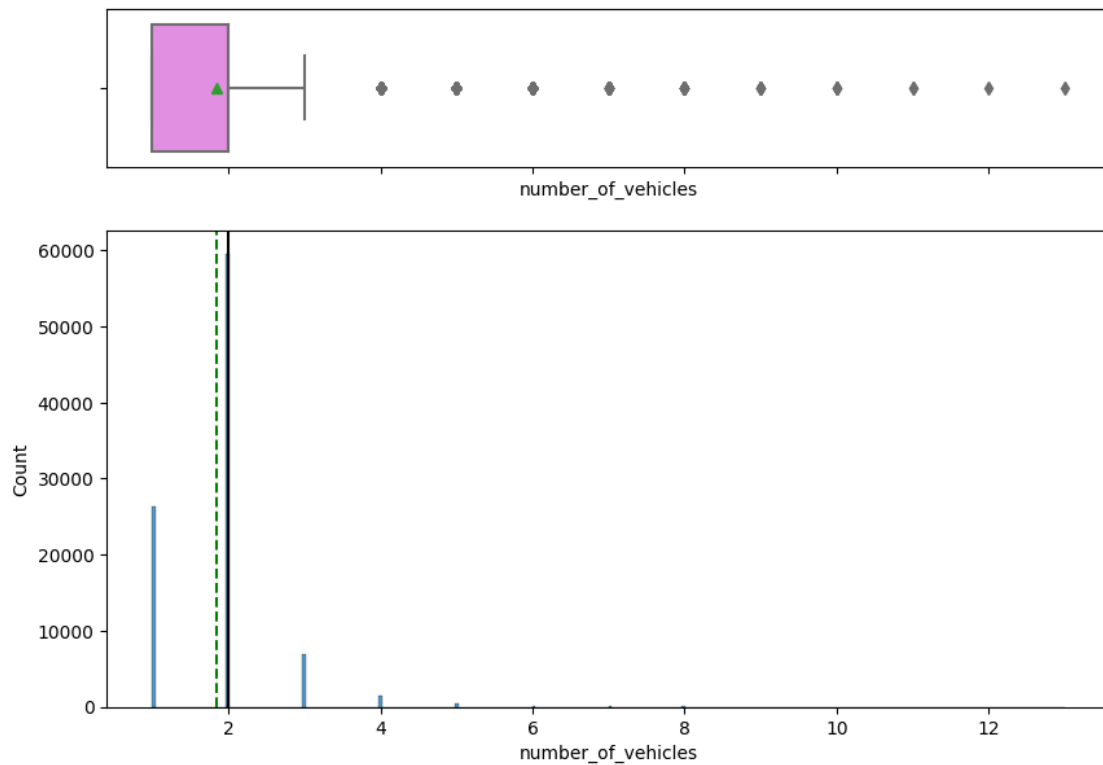
    data = data, x = feature, kde = kde, ax = ax_hist2
)
ax_hist2.axvline(
    data[feature].mean(), color = "green", linestyle = "--"
)
ax_hist2.axvline(
    data[feature].median(), color = "black", linestyle = "-"
)

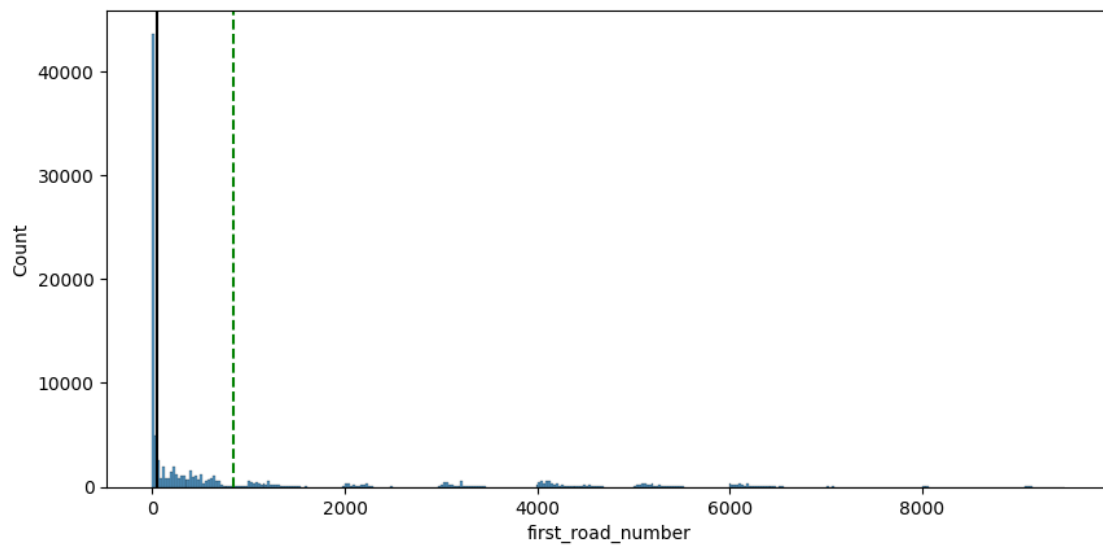
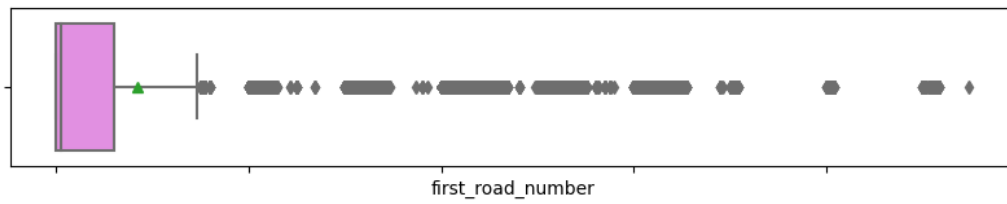
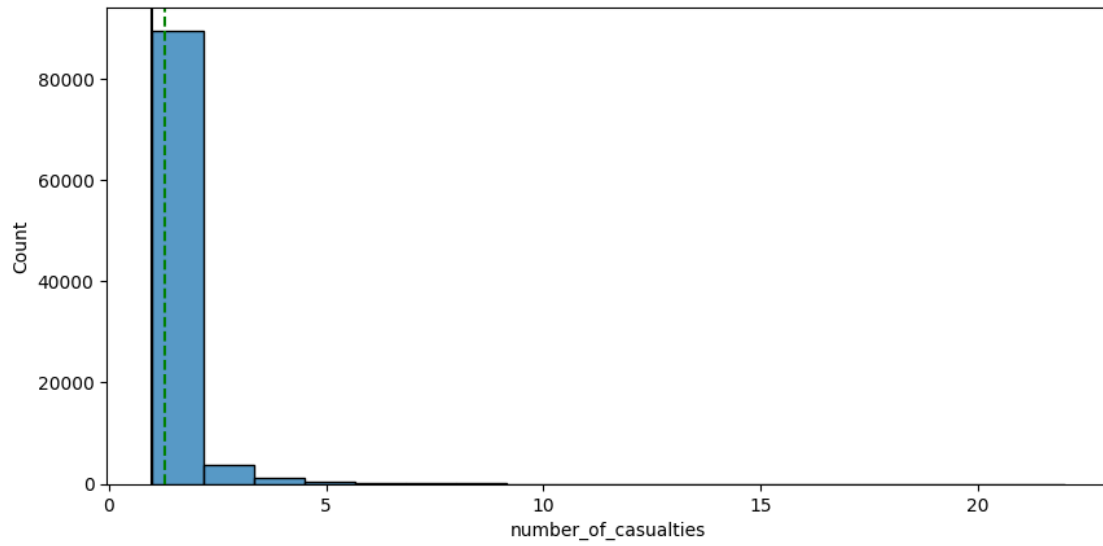
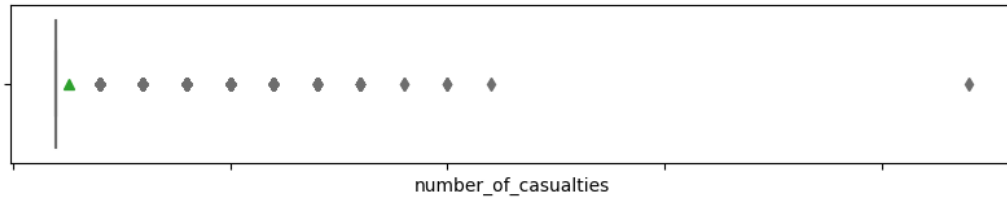
```

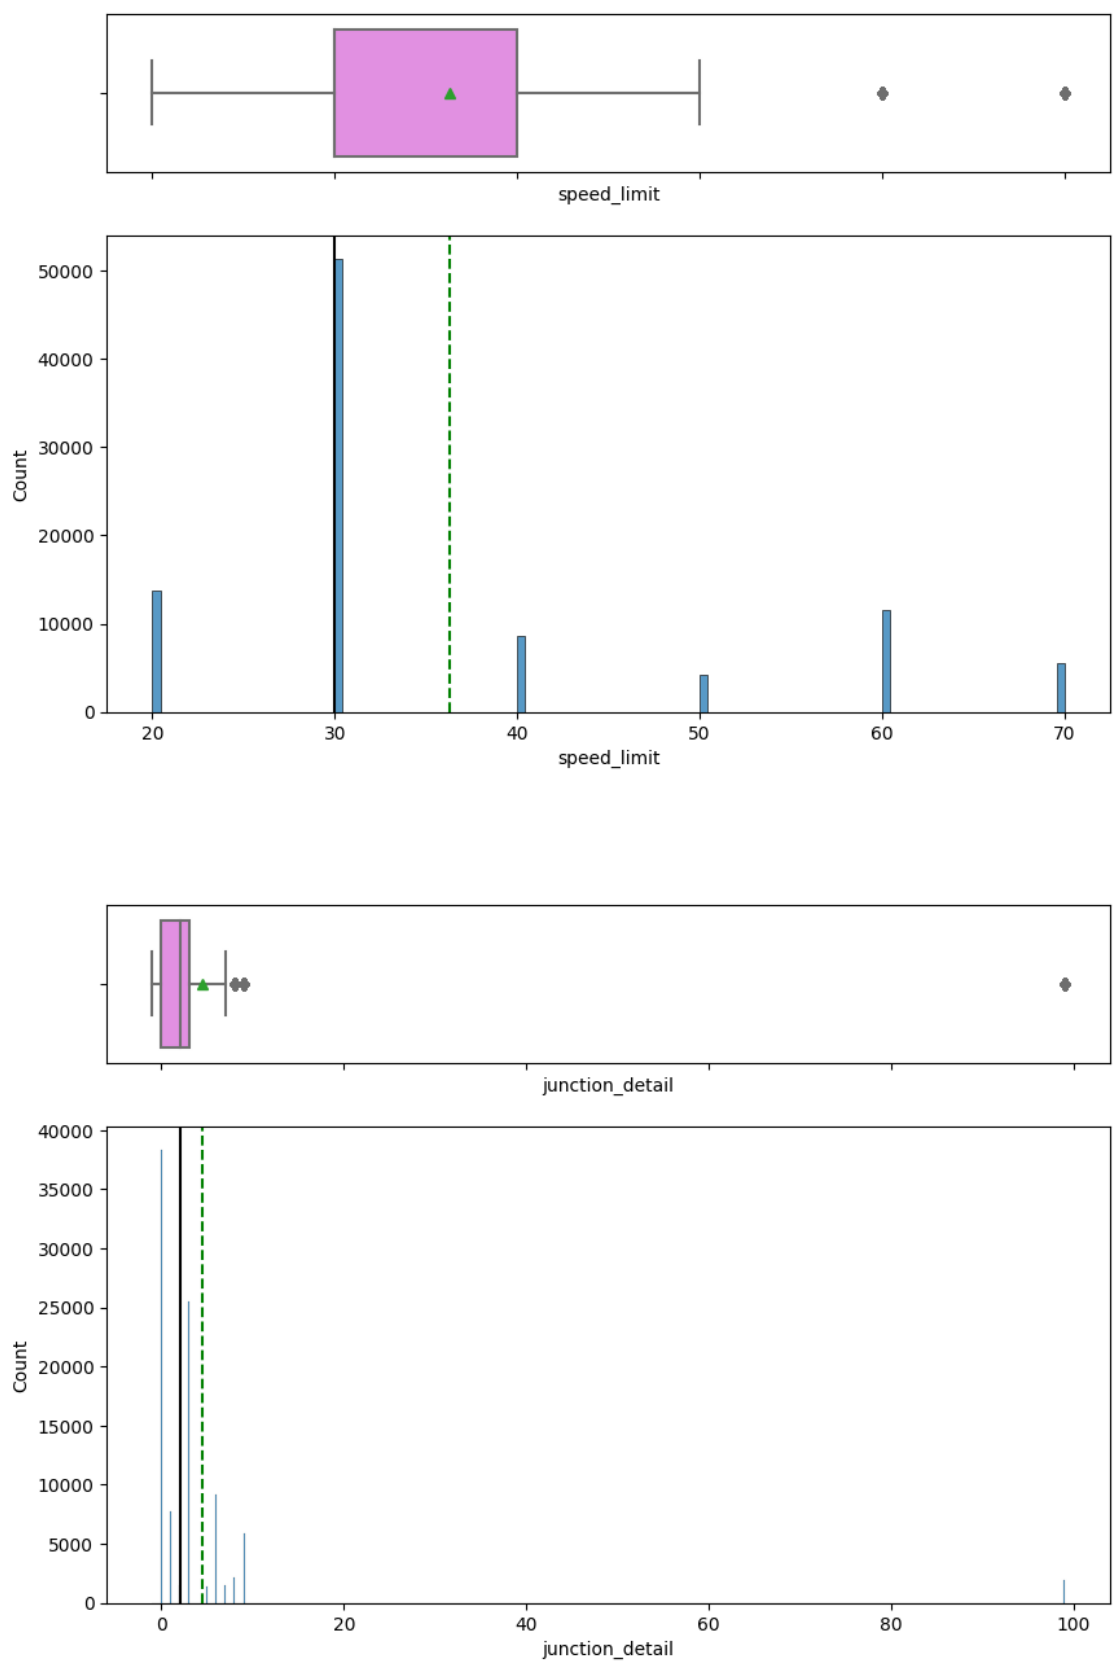
```

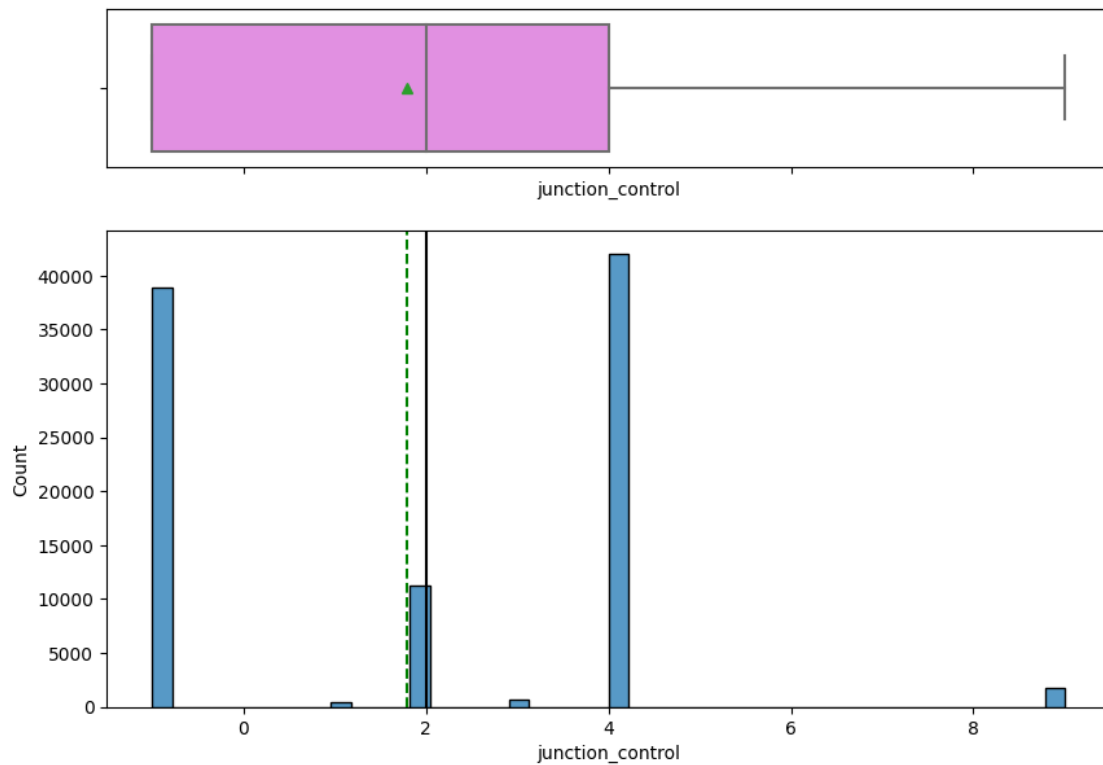
[44]: for i in num_cols:
    histogram_boxplot(data, i)

```









[45]: *# Function to create labeled barplots*

```
def labeled_barplot(data, feature, perc = False, n = None):
```

```

    total = len(data[feature])
    count = data[feature].nunique()
    if n is None:
        plt.figure(figsize = (count + 1, 5))
    else:
        plt.figure(figsize = (n + 1, 5))

    plt.xticks(rotation = 90, fontsize = 15)
    ax = sns.countplot(
        data = data,
        x = feature,
        palette = "Paired",
        order = data[feature].value_counts().index[:n].sort_values(),
    )

```

```

for p in ax.patches:
    if perc == True:
        label = "{:.1f}%".format(
            100 * p.get_height() / total
        )
    else:
        label = p.get_height()

    x = p.get_x() + p.get_width() / 2
    y = p.get_height()

    ax.annotate(
        label,
        (x, y),
        ha = "center",
        va = "center",
        size = 12,
        xytext = (0, 5),
        textcoords = "offset points",
    )

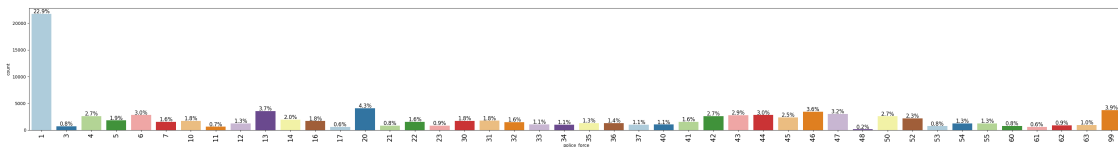
plt.show()

```

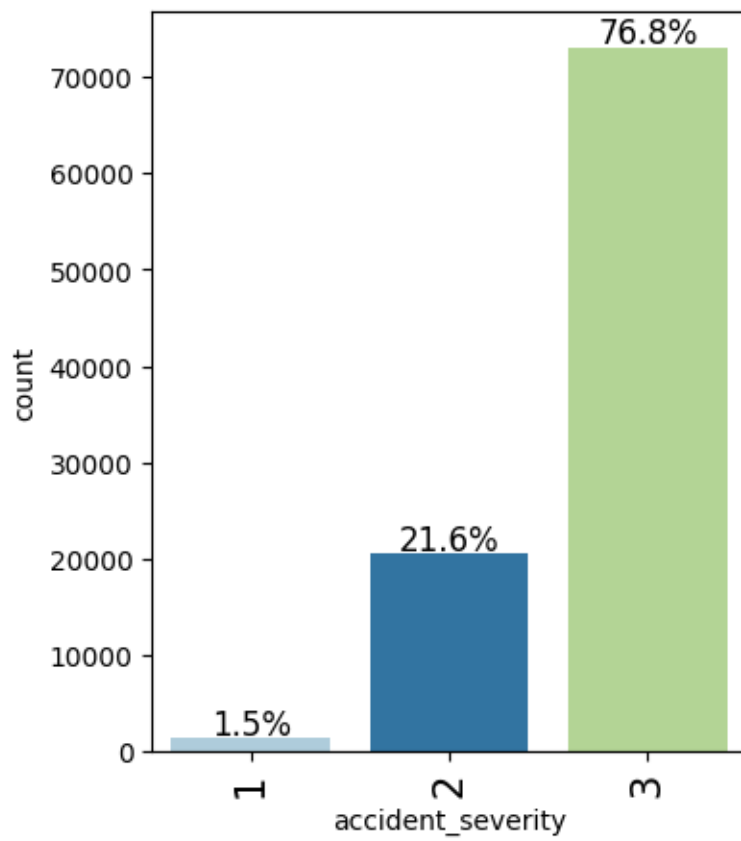
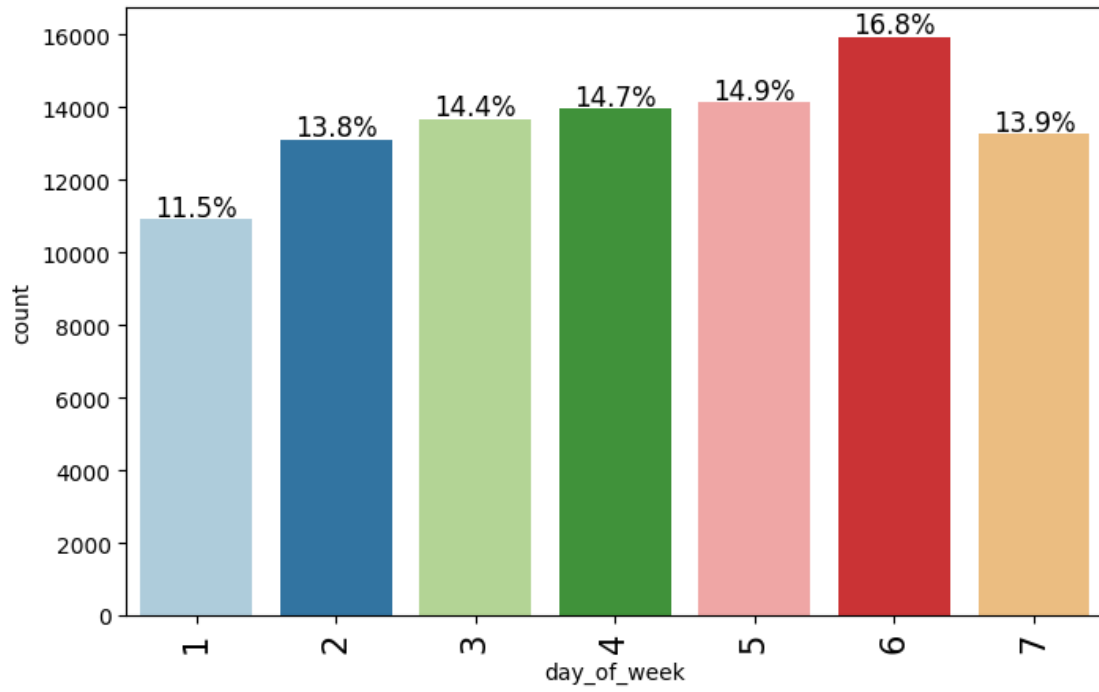
```

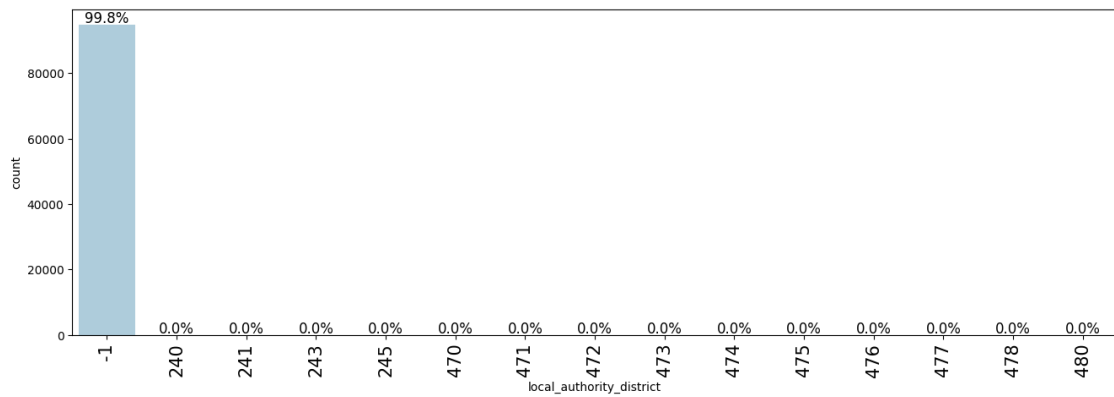
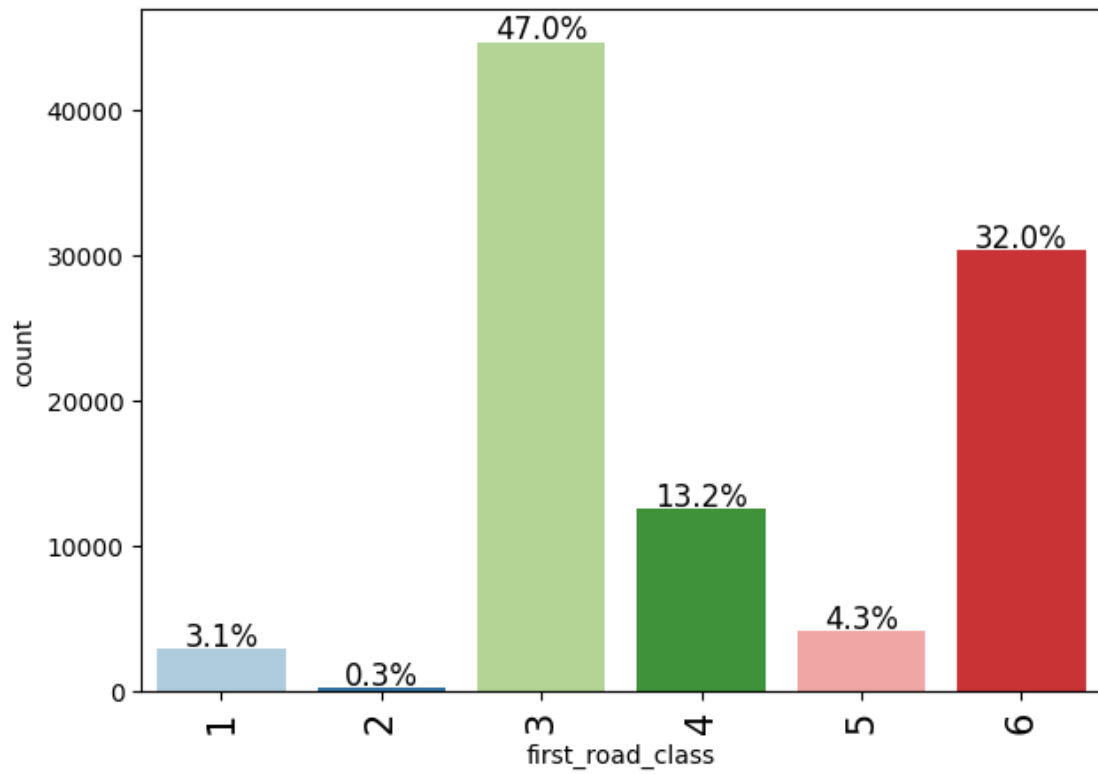
[46]: for i in cat_cols:
        labeled_barplot(data, i, perc = True)

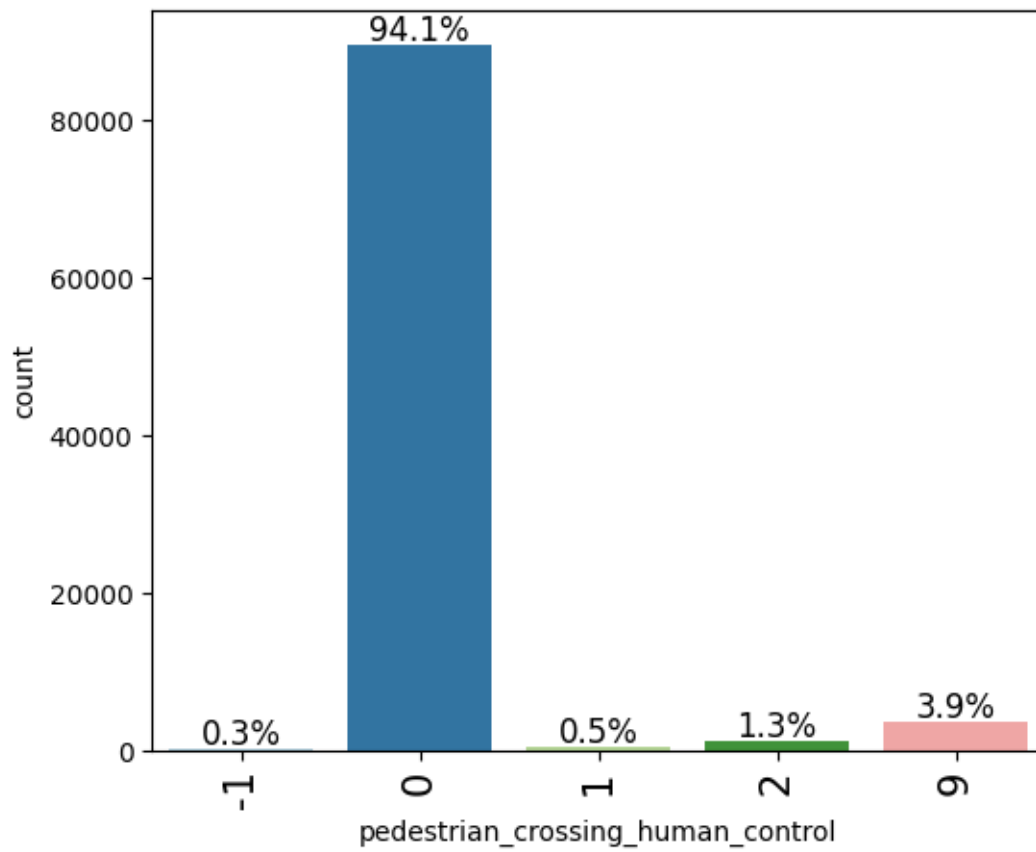
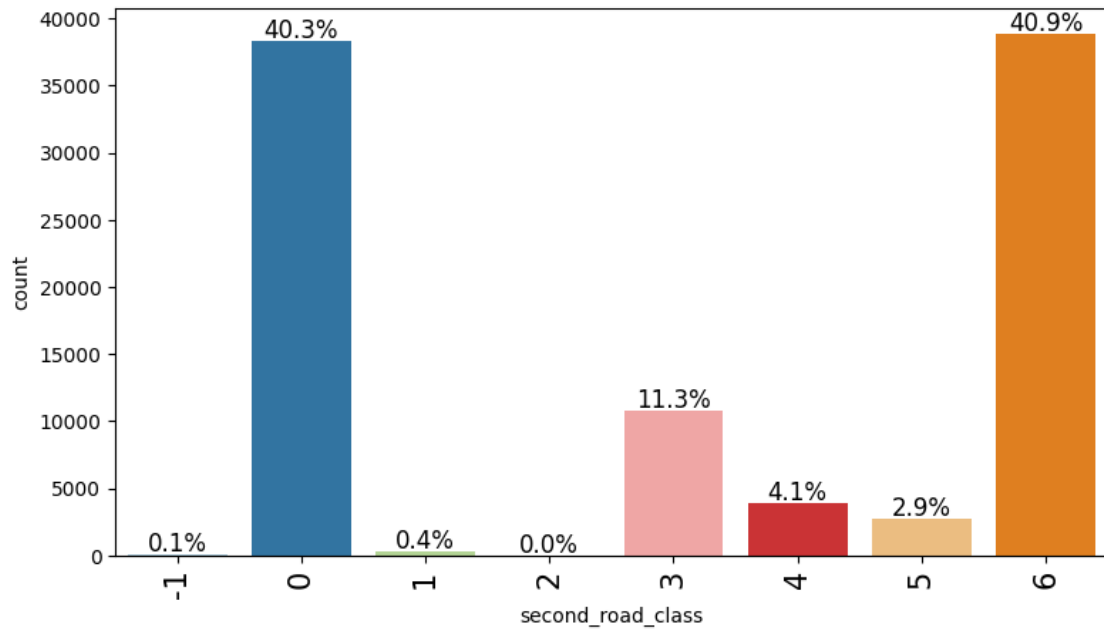
```

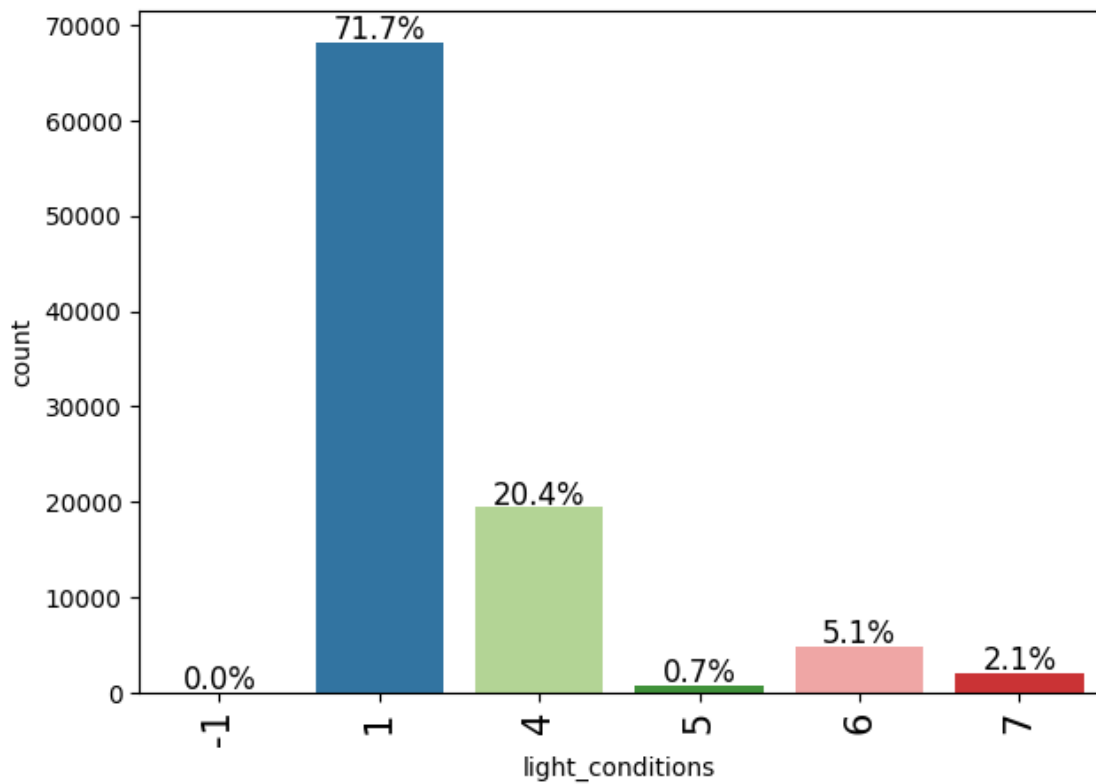
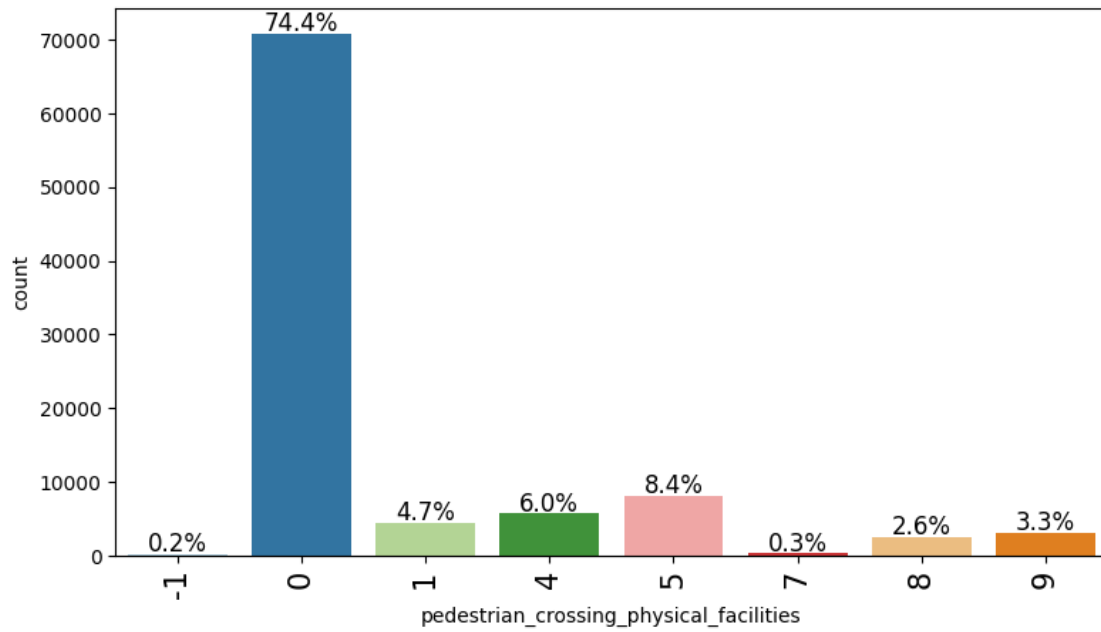


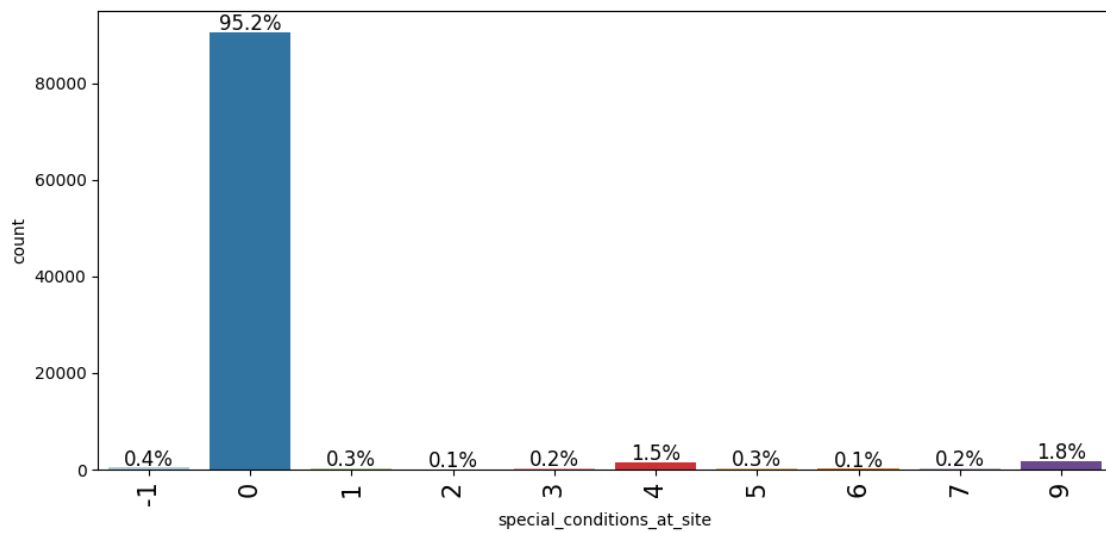
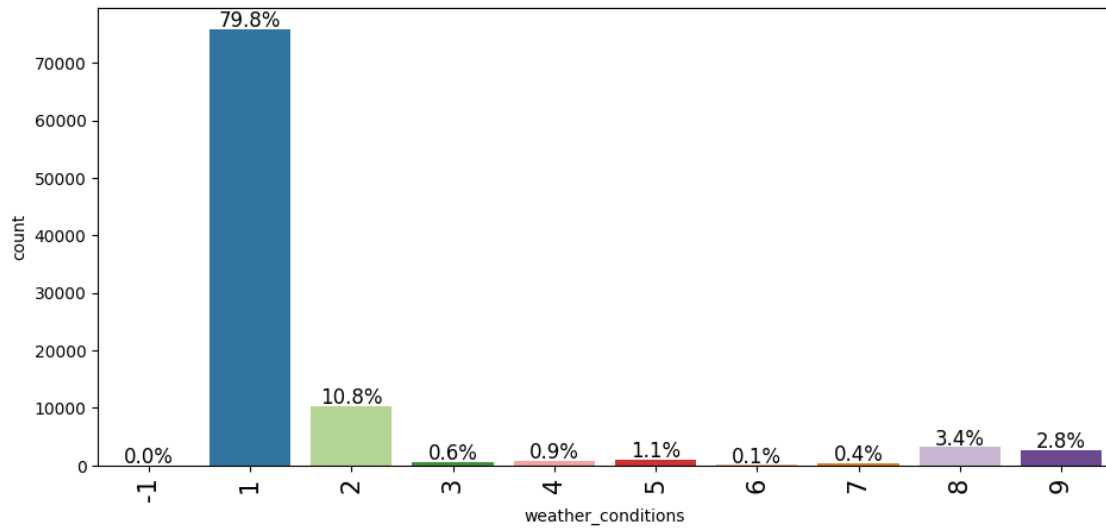


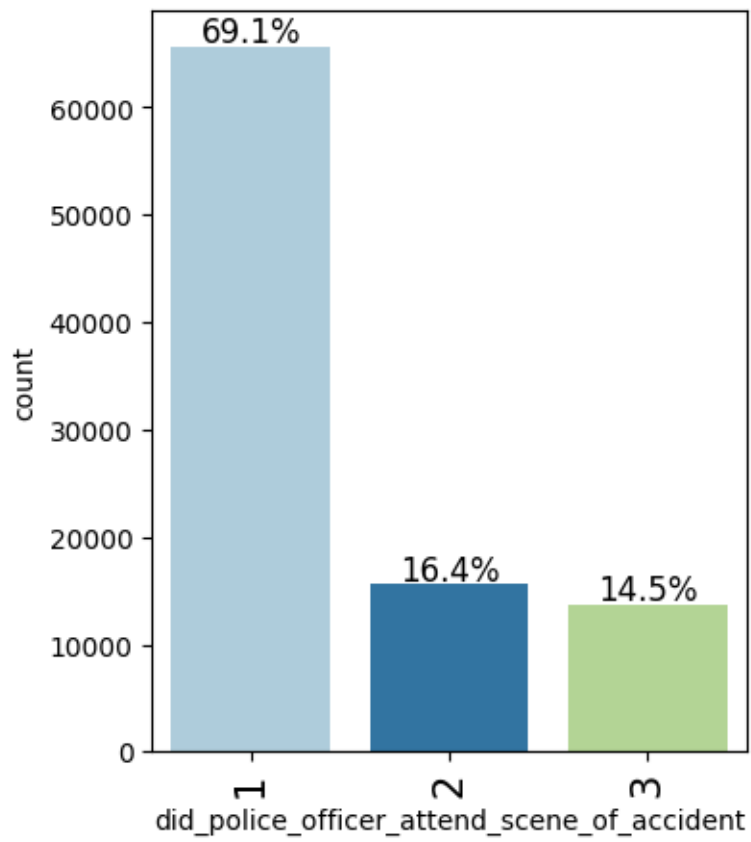
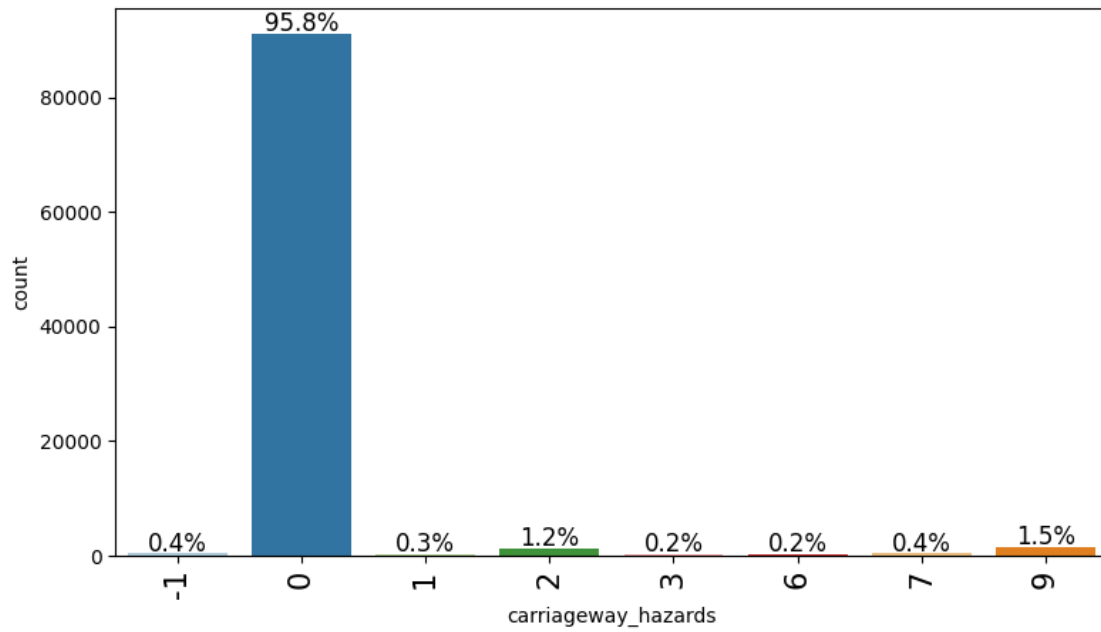


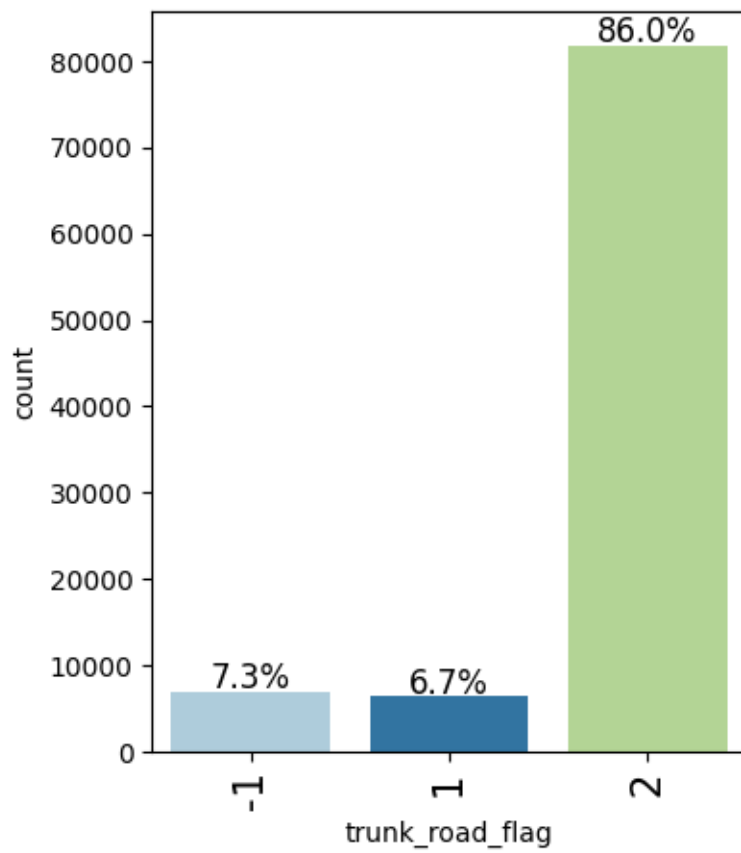
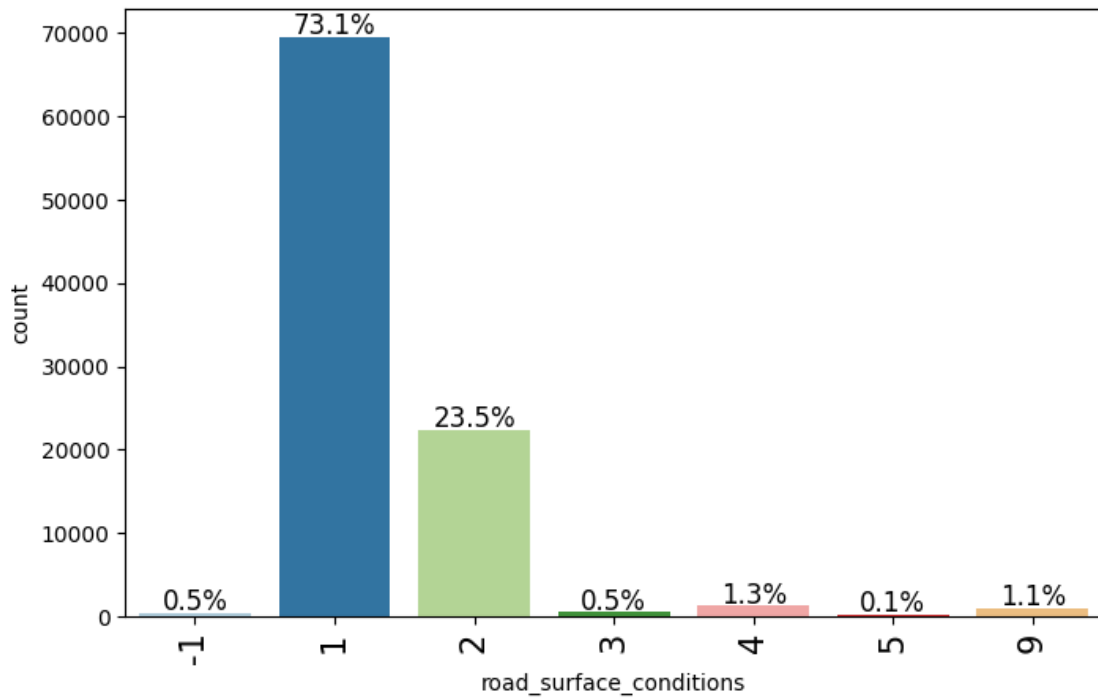


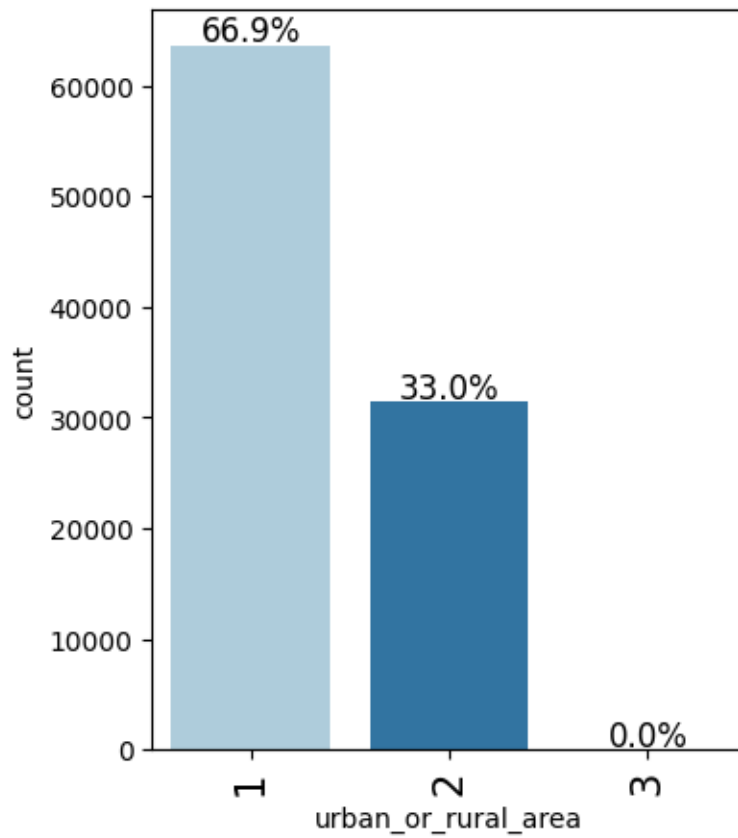






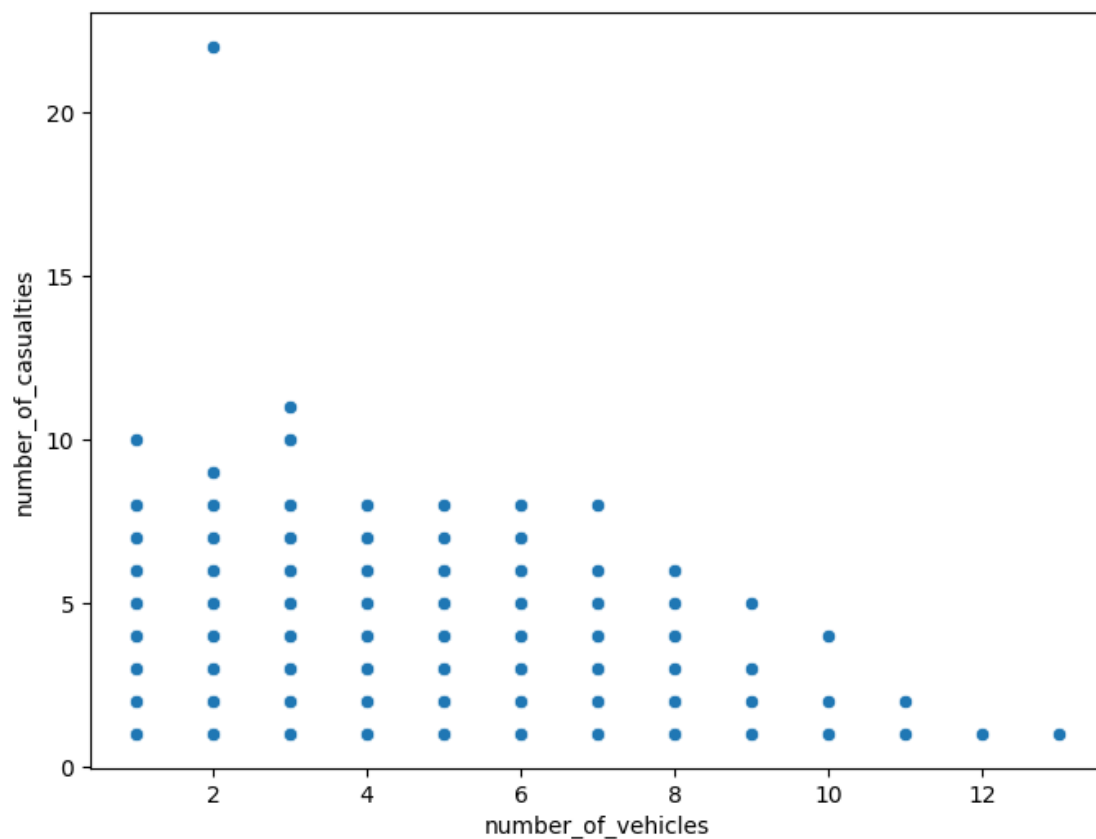






```
[48]: plt.figure(figsize = [8, 6])  
sns.scatterplot(x = data.number_of_vehicles, y = data.number_of_casualties)  
plt.show()
```





```
[47]: data.head()
```

```
[47]:  police_force accident_severity number_of_vehicles number_of_casualties \
0      1      1      3      3      1
1      1      1      2      2      3
2      1      1      2      2      4
3      1      1      1      1      1
4      1      1      3      4      1

   day_of_week local_authority_district first_road_class first_road_number \
0      6      -1      6      0
1      6      -1      3      1203
2      6      -1      4      272
3      6      -1      3      3220
4      6      -1      5      0

   road_type speed_limit junction_detail junction_control \
0      6      30      9      4
1      3      30      7      2
2      6      30      9      2
```

3	2	30	9	4
4	6	20	3	4

	second_road_class	second_road_number	pedestrian_crossing_human_control	\
0	6	0	0	
1	3	1204	0	
2	5	0	0	
3	6	0	0	
4	6	0	0	

	pedestrian_crossing_physical_facilities	light_conditions	weather_conditions	\
0	0	4	7	
1	5	4	1	
2	5	4	1	
3	4	4	1	
4	0	4	1	

	road_surface_conditions	special_conditions_at_site	carriageway_hazards	\
0	4	1	0	
1	1	0	0	
2	1	0	0	
3	1	0	0	
4	1	0	0	

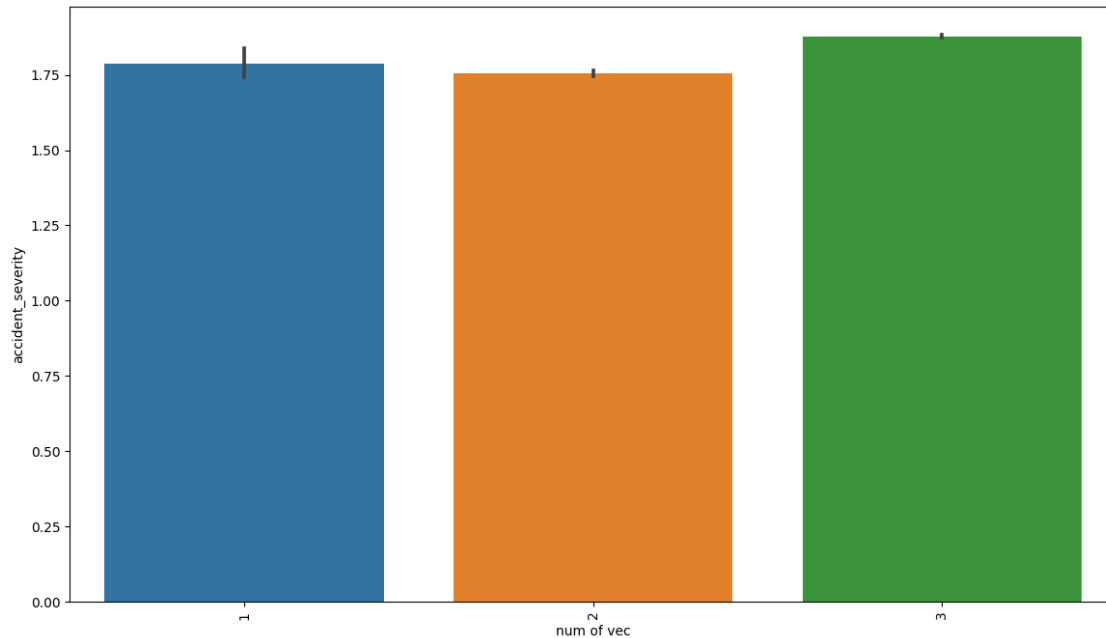
  

	urban_or_rural_area	did_police_officer_attend_scene_of_accident	\
0	1	1	
1	1	1	
2	1	1	
3	1	1	
4	1	1	

	trunk_road_flag
0	2
1	2
2	2
3	2
4	2

```
[50]: df_revenue1 = data.groupby(["number_of_vehicles"], as_index = False)[
        "accident_severity"
    ].sum()
plt.figure(figsize = [14, 8])
plt.xticks(rotation = 90)
a = sns.barplot(x = data.accident_severity, y = data.number_of_vehicles)
a.set_xlabel("num of vec")
a.set_ylabel("accident_severity")
plt.show()
```



```
[54]: # Creating the list of columns for which we need to create the dummy variables
to_get_dummies_for = ['police_force',
    'day_of_week',
    'first_road_class',
    'local_authority_district',
    'second_road_class',
    'pedestrian_crossing_human_control',
    'pedestrian_crossing_physical_facilities',
    'light_conditions',
    'weather_conditions',
    'special_conditions_at_site',
    'carriageway_hazards',
    'did_police_officer_attend_scene_of_accident',
    'road_surface_conditions',
    'trunk_road_flag',
    'urban_or_rural_area']

# Creating dummy variables
data = pd.get_dummies(data = data, columns = to_get_dummies_for, drop_first = 
    True)
```

```
[56]: # Separating the target variable and other variables
X = data.drop(columns = ['accident_severity'])
y = data.accident_severity
```

```
[57]: # Splitting the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
↳ random_state = 1, stratify = y)
```

Model Building

```
[58]: # reference for function is used fro google.
def model_performance_classification_sklearn(model, predictors, target):

    pred = model.predict(predictors)

    acc = accuracy_score(target, pred)
    recall = recall_score(target, pred, average = "weighted")
    precision = precision_score(target, pred, average = 'weighted')
    f1 = f1_score(target, pred, average = 'weighted')

    df_perf = pd.DataFrame(
        {"Accuracy": acc, "Recall": recall, "Precision": precision, "F1": f1,},
        index=[0],
    )

    return df_perf
```

Training models on original data

```
[65]: models = []

models.append(("Logistic regression", LogisticRegression(random_state=1)))
models.append(("dtree", DecisionTreeClassifier(random_state=1)))
models.append(("Random forest", RandomForestClassifier(random_state=1)))

print("\n" "Training Performance:" "\n")
for name, model in models:
    model.fit(X_train, y_train)
    scores_train = recall_score(y_train, model.predict(X_train), average =
↳ "weighted")
    print("{}: {}".format(name, scores_train))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train, y_train)
    scores_val = recall_score(y_test, model.predict(X_test), average =
↳ "weighted")
    print("{}: {}".format(name, scores_val))
```

#### Training Performance:

Logistic regression: 0.7685047417226506  
dtree: 0.9869546267490269  
Random forest: 0.9869395975171709

#### Validation Performance:

Logistic regression: 0.7684808528545378  
dtree: 0.6419553934633189  
Random forest: 0.7416888764202553

- The training performance of the decision tree and random forest models is very high, with accuracy scores of 0.9869 and 0.9869 respectively. This indicates that these models are likely overfitting the training data, as they are performing almost perfectly on the training set but not as well on the validation set.
- In contrast, the training performance of the logistic regression model is lower, with an accuracy score of 0.7685. This suggests that the logistic regression model may be less prone to overfitting, and may be a more robust model in this case. The validation performance of the logistic regression model is similar to its training performance, with an accuracy score of 0.7685. This suggests that the model is likely performing consistently across both the training and validation sets, and may be a good candidate for deployment in practice.
- The validation performance of the decision tree model is much lower than its training performance, with an accuracy score of 0.6420. This indicates that the model is overfitting the training data, and is not generalizing well to new data.
- The validation performance of the random forest model is also lower than its training performance, but not as much as the decision tree model. Its accuracy score on the validation set is 0.7417, which is still a relatively good performance, but may indicate some overfitting.

```
[63]: print("Before Oversampling, '1': {}".format(sum(y_train == 1)))  
      print("Before Oversampling, '2': {}".format(sum(y_train == 2)))  
      print("Before Oversampling, '3': {}".format(sum(y_train == 3)))
```

Before Oversampling, '1': 1030  
Before Oversampling, '2': 14373  
Before Oversampling, '3': 51134

```
[64]: sm = SMOTE(  
      sampling_strategy="auto", k_neighbors=3, random_state=1  
      ) # Synthetic Minority Over Sampling Technique  
      X_train_over, y_train_over = sm.fit_resample(X_train, y_train)
```

```
[66]: models = []  
  
models.append(("Logistic regression", LogisticRegression(random_state=1)))  
models.append(("dtree", DecisionTreeClassifier(random_state=1)))  
models.append(("Random forest", RandomForestClassifier(random_state=1)))
```

```

print("\n" "Training Performance:" "\n")
for name, model in models:
    model.fit(X_train_over, y_train_over)
    scores_train = recall_score(y_train_over, model.predict(X_train_over),
    ↪average = "weighted")
    print("{}: {}".format(name, scores_train))

print("\n" "Validation Performance:" "\n")

for name, model in models:
    model.fit(X_train, y_train)
    scores_val = recall_score(y_test, model.predict(X_test), average =
    ↪"weighted")
    print("{}: {}".format(name, scores_val))

```

Training Performance:

Logistic regression: 0.4662977014641269  
 dtree: 0.9880118903273751  
 Random forest: 0.9880053715075423

Validation Performance:

Logistic regression: 0.7684808528545378  
 dtree: 0.6419553934633189  
 Random forest: 0.7416888764202553

- The training performance of the decision tree and random forest models is very high (close to 1.0), while the training performance of the logistic regression model is relatively low (around 0.47). This suggests that the decision tree and random forest models may be overfitting the training data, while the logistic regression model may not be complex enough to capture the patterns in the data.
- The validation performance of the logistic regression model is better than the other models (0.77 vs. 0.64 and 0.74 for the decision tree and random forest, respectively). This suggests that the logistic regression model may be better at generalizing to new, unseen data than the other models.
- The validation performance of the decision tree model is the worst among the three models. This is consistent with the observation that the decision tree model may be overfitting the training data.
- The random forest model has a relatively high training performance but a lower validation performance than the logistic regression model. This suggests that the random forest model may be overfitting the training data to some extent, but not as severely as the decision tree model.
- It may be a good idea to try tuning the hyperparameters of the decision tree and random

forest models using techniques like grid search or random search to see if their performance can be improved. It may also be worth exploring other classification algorithms to see if they can achieve better performance on this dataset.

Tuning random forest on the original data

```
[68]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
import pandas as pd
import numpy as np

# Define the parameter grid for Grid Search
param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [10, 20],
    'max_features': [3, 6],
    'min_samples_split': [2, 4],
    'min_samples_leaf': [1, 2]
}

# Define the Random Forest model
rf = RandomForestClassifier(random_state=42)

# Perform Grid Search to find the best parameters
grid_search = GridSearchCV(rf, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(X_train, y_train)
```

```
[ ]: ## Print the best parameters and accuracy score
# print("Best parameters: ", grid_search.best_params_)
# print("Best accuracy score: ", grid_search.best_score_)

# Make predictions using the best model
best_rf = grid_search.best_estimator_
y_pred = best_rf.predict(X_train)

## Evaluate the performance of the best model
# accuracy = accuracy_score(y_test, y_pred)
# print("Accuracy score: ", accuracy)

# scores_train = recall_score(y_train, model.predict(X_train), average = "weighted")
# print("{Model}: {}".format(name, scores_train))
```

Feature importance

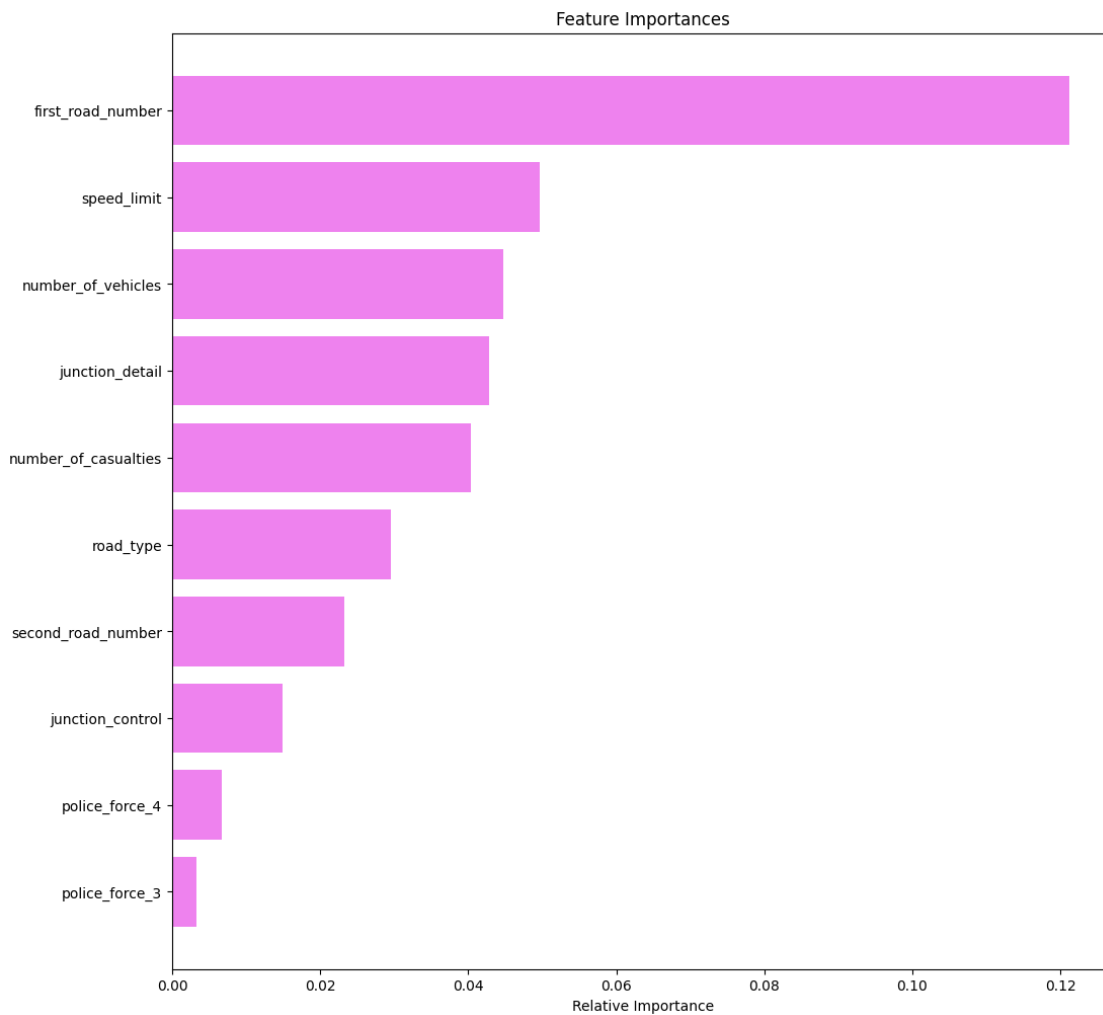
```
[74]: best_rf = models[2][1]
feature_names = X_train.columns
```

```

importances = best_rf.feature_importances_[:10]
indices = np.argsort(importances)

plt.figure(figsize=(12, 12))
plt.title("Feature Importances")
plt.barh(range(len(indices)), importances[indices], color="violet",
         align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()

```



- we can also observe the top 10 features

### 0.1.1 Summary/Conclusion

EDA:



We performed exploratory data analysis (EDA) on the accident severity dataset and found that there are 24 features in total, including factors such as the number of vehicles involved, the weather conditions, and whether a police officer attended the scene. We also observed that some features are categorical and others are numerical, and that there are some missing values in the dataset.

### **Model Development:**

We then built three different machine learning models - logistic regression, decision tree, and random forest - to predict accident severity based on the features in the dataset. We trained these models on a training set and evaluated their performance on a validation set.

### **Results:**

Our analysis showed that the decision tree and random forest models had very high accuracy scores on the training set, but lower scores on the validation set. This suggests that these models are overfitting the training data and may not generalize well to new data. The logistic regression model, on the other hand, had a lower accuracy score on the training set, but a similar score on the validation set, indicating that it may be a more robust model for predicting accident severity.

### **Conclusion:**

Based on our analysis, we recommend using a logistic regression model to predict accident severity, as it appears to be more robust and less prone to overfitting. We also recommend further investigation into the most important features for predicting accident severity, as this information could be useful for developing targeted interventions to reduce the number and severity of accidents. Overall, our analysis demonstrates the potential of machine learning techniques for improving road safety and reducing the human and economic costs of accidents.

[ ]: